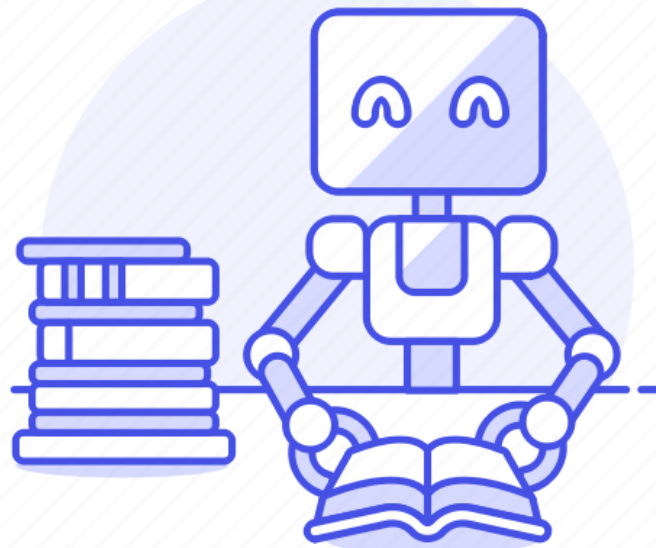# L2 Introduction to Machine Learning

Zonghua Gu 2022

ML Taxonomy

- ## Supervised Learning:
  - The system is presented with example inputs and their desired outputs, given by a "teacher", and the goal is to learn a general rule that maps inputs to outputs.
    - Classification (cat or dog?)
    - Regression (housing price next year?)
- ## Unsupervised Learning:
  - No labels are given to the learning algorithm, leaving it on its own to find structure in its input. Unsupervised learning can be a goal in itself (discovering hidden patterns in data) or a means towards an end (feature learning).
    - Parametric UL (e.g., Gaussian Mixture Models)
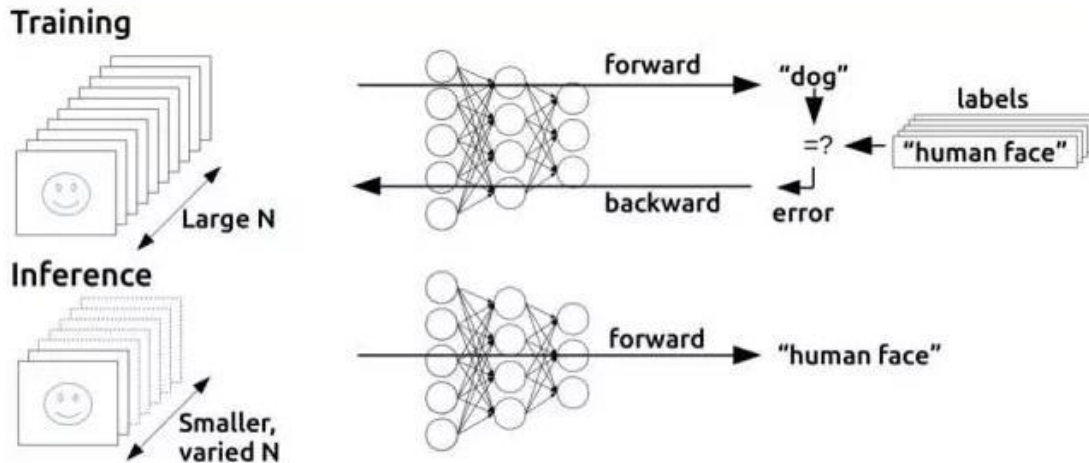   - Non-parametric UL

ML Taxonomy

- Reinforcement Learning:
  - An agent interacts with a dynamic environment in which it must perform a certain goal. The agent is provided feedback in terms of rewards and it tries to learn an optimal policy that maximizes its cumulative rewards.
  - Algorithms: Model-based; Model-free (Value-based, Policy-based)
  - Applications: Game playing (AlphaGo); Robotics; AD…

state S_t reward R_t

R_{t+1}

S_{t+1}

Agent

Environment
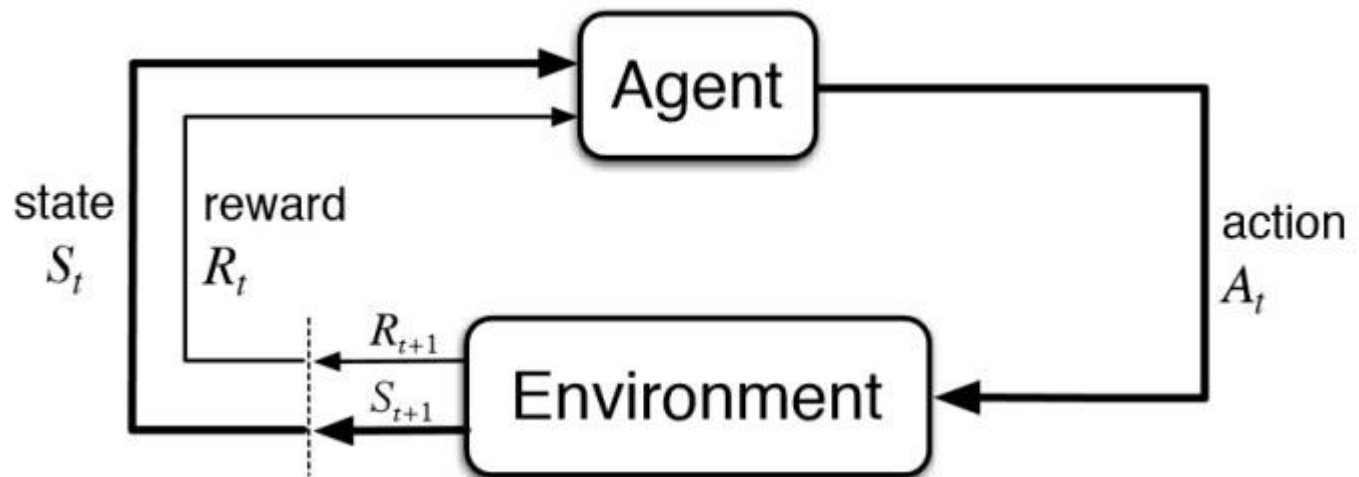
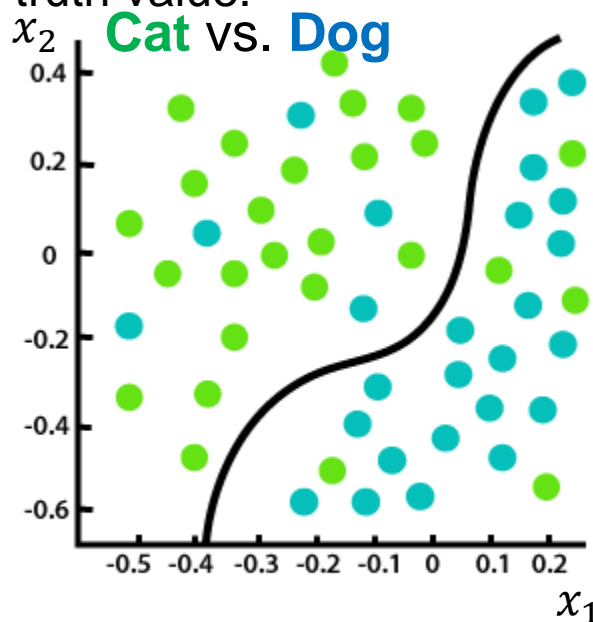action A_t

Training vs. Inference

- Training: millions of iterations of forward pass + back propagation to adjust model params (e.g., NN weights); requires large CPU/GPU clusters and days/weeks of training time
- Inference (also called prediction): a single forward pass; can be run on edge devices

Training

Inference

Large N

Smaller, varied N

forward
"dog"

backward
error

=?

labels
"human face"

forward → "human face"

```
def train(train_images, train_labels):
    # build a model for images -> labels...
    return model


def predict(model, test_images):
    # predict test_labels using the model...
    return test_labels
```

Supervised Learning: Classification and Regression

- Classification is used to predict/classify discrete labels such as Male or Female, True or False, Spam or Not Spam, etc.
- Regression is used to predict continuous values such as price, salary, age, etc.
- Both are Supervised Learning algorithms that require ground-truth values as labels.
- Both need loss functions to measure how the predicted value differs from ground-truth value.

**Classification**　　　　**Regression**

5
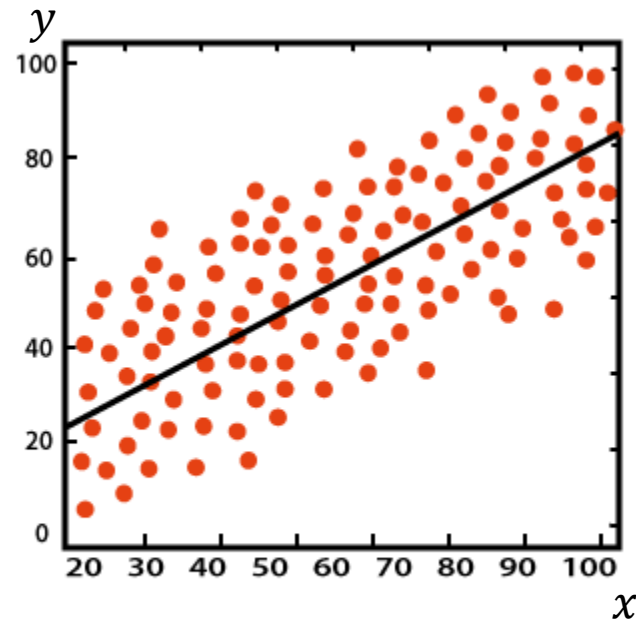
A Neuron and its Activation Function

- The activation function is a nonlinear monotonic function that acts like a "gate": the output is larger for larger input activation
  - Perceptron $y = \sigma(z) = \text{step}(wx + b)$ (activation function $f = $ step function, shown below)
  - Linear Regression if $y = z = wx + b$ (activation function $f = $ identity function)
  - Logistic Regression if $y = \sigma(z) = \sigma(wx + b)$ (activation function $f = $ sigmoid function)
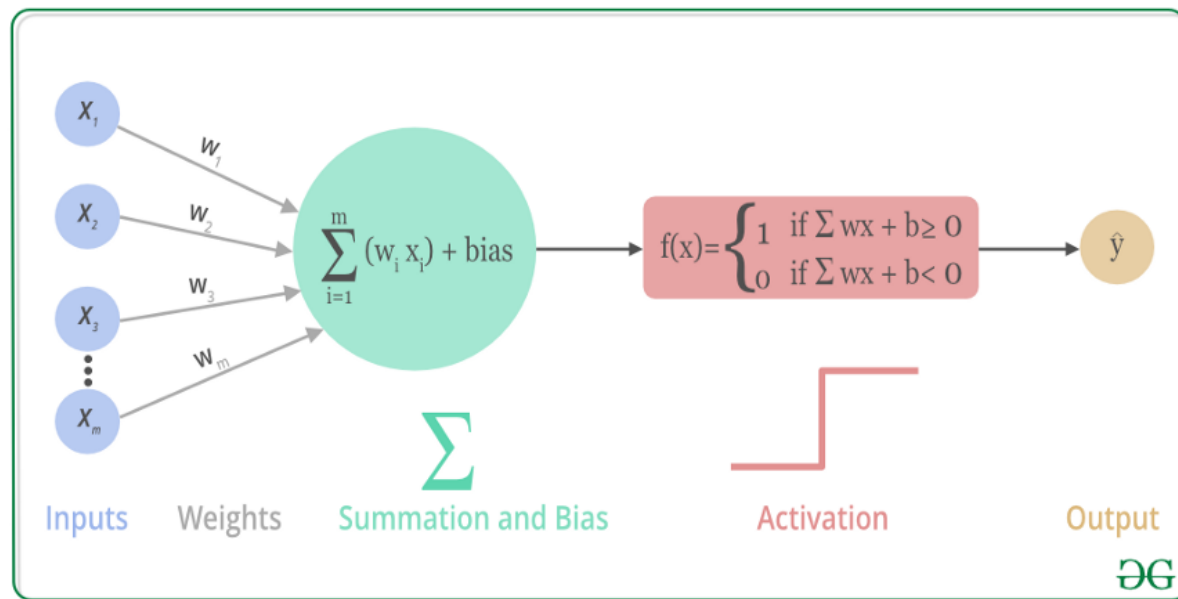
Linear Regression for Regression

- Function approximation $y = wx + b$, with learnable parameters $\theta = \{w, b\}$, where x, y, b are vectors, and w is a weight matrix
  - e.g., we want to predict price of a house based on its feature vector $\mathbf{x} = [x_1 \ x_2 \ x_3]^\mathbf{T}$, where x_1 is area in square meters (sqm), x_2 is location ranking (loc), x_3 is year of construction (yoc)
  - Predicted price $y = wx + b = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$
  - Fig shows an example for scalar x and y

Linear Regression example. y = 7x + N(0, 200)

Logistic Regression for Binary Classification

- Consider a binary classification problem: an input image x may be classified as a dog with probability $P(y = dog|x)$, a cat with probability $P(y = cat|x)$, with $P(y = dog|x) + P(y = cat|x) = 1.0$

- Logistic Regression: use sigmoid function $\sigma(z_i) = \frac{1}{1+e^{-z_i}}$ to map from the activation (also called the logit) to the output probability

- In addition to binary classification at the output layer, sigmoid may also be used as the non-linear activation function in the hidden layers of a NN

Class 1

classifier

Class 2
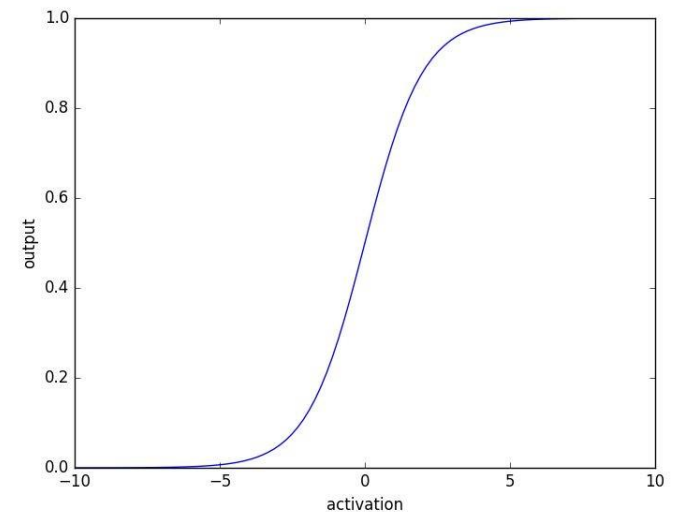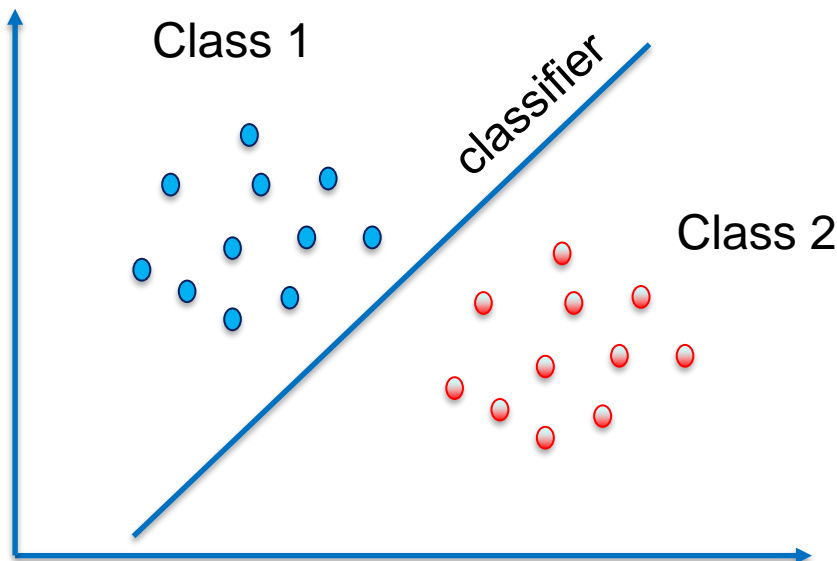
Common Activation Functions used in DL

Sigmoid

$\sigma(x) = \frac{1}{1+e^{-x}}$

tanh

$\tanh(x)$

ReLU

$\max(0, x)$
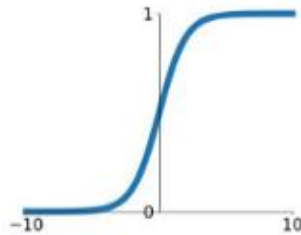
Leaky ReLU
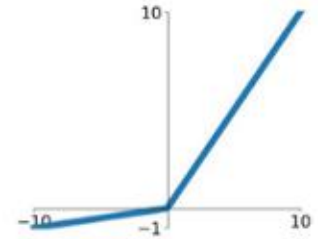
$\max(0.1x, x)$

Maxout

$\max(w_1^T x + b_1, w_2^T x + b_2)$
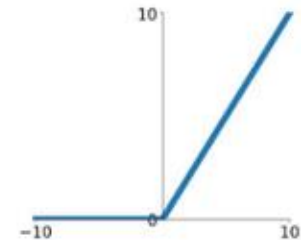
ELU

$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$
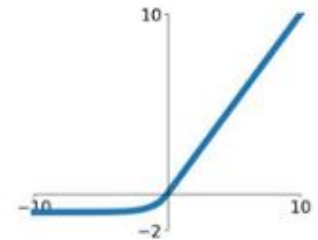
Deep Neural Networks

- We can stack many hidden layers to form a DNN if we have enough data and computing power to train it

- The high model capacity of DNN comes from non-linear mappings: hidden units must be followed by a non-linear activation function

  – Without non-linear activation functions, a DNN with many layers can be collapsed into an equivalent single-layer NN

Fully-Connected NNs

- Number of params at the i-th layer is $(N_{i-1} + 1) * N_i$, where N_i is the number of neurons at the i-th layer. Can grow very large
  - (We will discuss CNNs with much fewer params in the next lecture)

Slide Credit: Hugo Laroche NN course

A 3-layer NN

Example: Two-Layer Fully-Connected NN for Solving XOR

- The NN consists of one input, one hidden, and one output layer, with sigmoid activations

x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0

2 hidden units with sigmoid activations

output unit with sigmoid activation
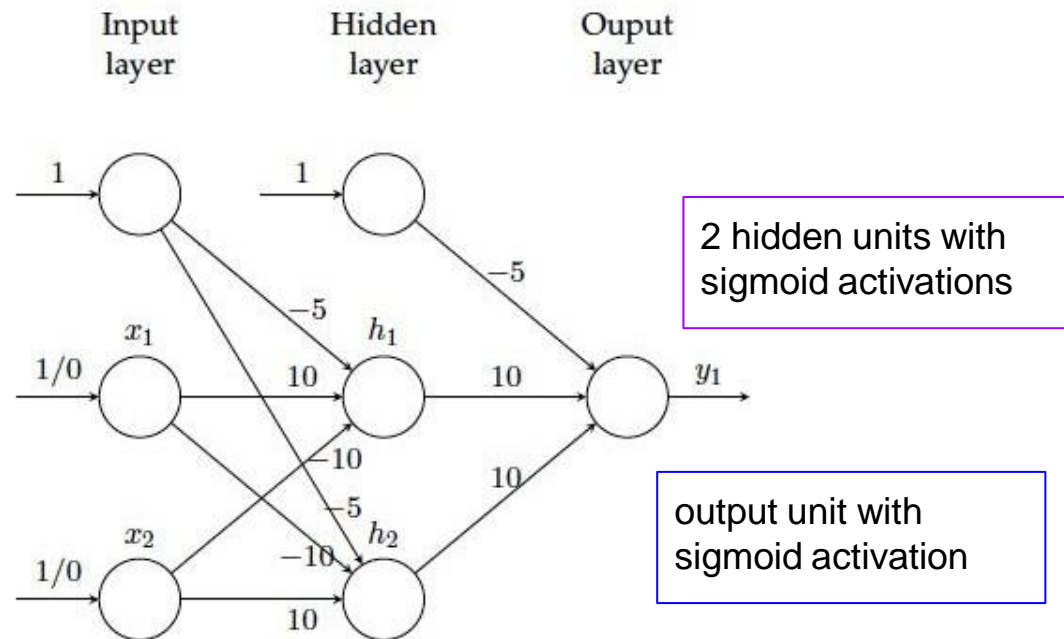
Setting # Layers and Their Sizes

- An example illustrating adding more hidden neurons increases model capacity and reduces training error

- But too many layers and neurons may lead to overfitting

3 hidden neurons	6 hidden neurons	20 hidden neurons
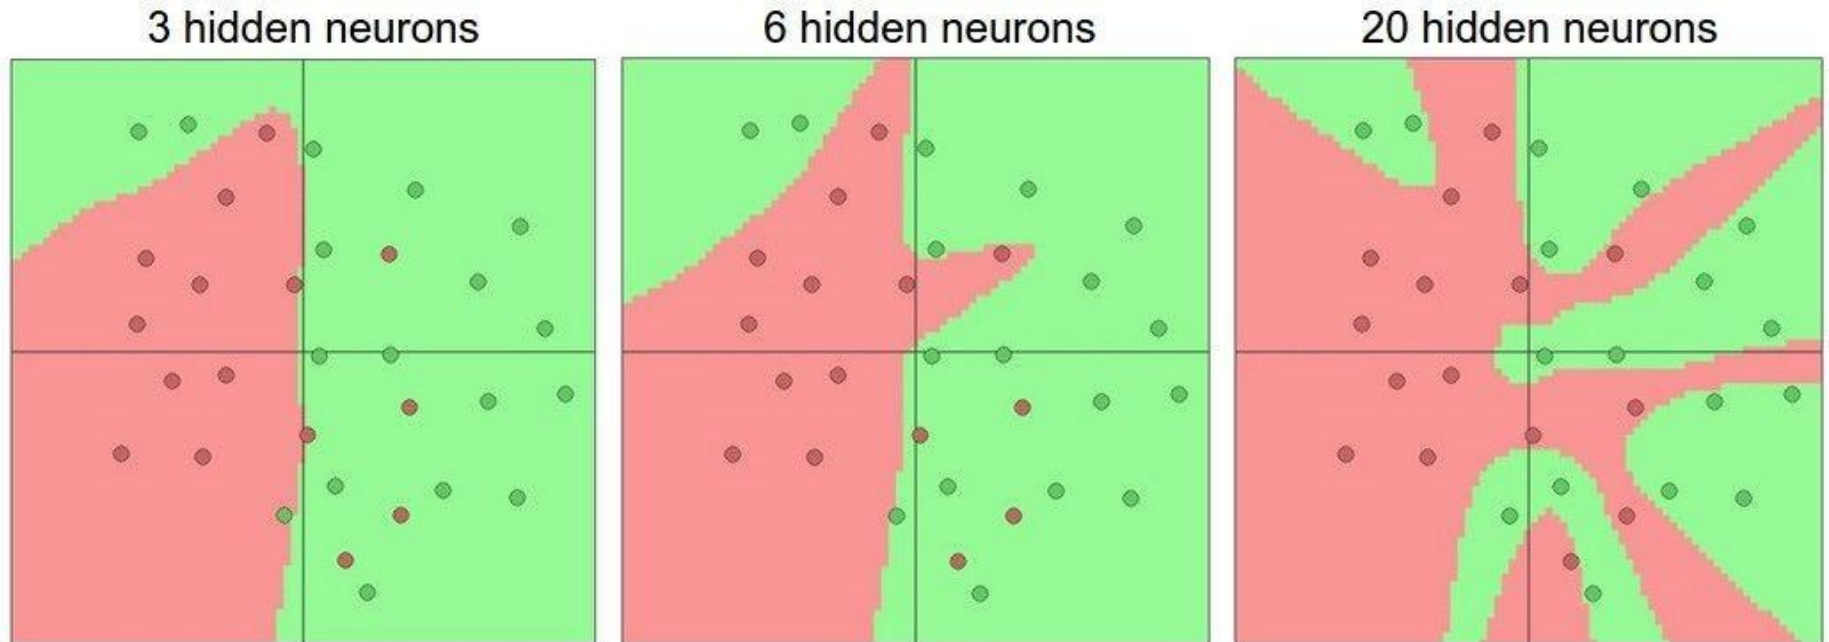
Loss Functions

- Classification
  - Cross-Entropy Loss, Log Loss, Focal Loss, Exponential Loss, Hinge Loss…

- Regression
  - MSE (Mean Squared Error)/L2 Loss/Quadratic Loss, MAE (Mean Absolute Error)/L1 Loss, Huber Loss, Log Cosh Loss, Quantile Loss…

NN for Multi-Class Classification

- Consider a NN defining the model $h_\theta: \mathcal{X} \to \mathbb{R}^k$, as the mapping from input x to output $h_\theta(x)$, a k-dim vector of logits, where k is the number of classes
  - θ is the set of params (weights and biases)
  - y is the correct label for input x
  - Note that h_θ does not include the last SoftMax layer

- e.g., a 3-layer NN consisting of 2 layers with ReLU activation functions and a last linear layer is
  - $h_\theta(x) = W_3 \max(0, W_2 \max(0, W_1 x + b_1) + b_2) + b_3$

Cross-Entropy Loss for Multi-Class Classification

- The SoftMax operator $\sigma: \mathbb{R}^k \to \mathbb{R}^k$ computes a vector of predicted probabilities $\sigma(z): \mathbb{R}^k$ from a vector of logits $z: \mathbb{R}^k$ in the last hidden layer (the penultimate layer), where k is the number of classes:

  - $\sigma(z)_i = \dfrac{\exp(z_i)}{\sum_{j=1}^{k} \exp(z_j)}$

- The loss function is defined as the negative log likelihood of the predicted probability corresponding to the correct label y:

  - $\text{Loss}(x, y; \theta) = -\log \sigma\big(h_\theta(x)\big)_y = -\log\left(\dfrac{\exp\big(h_\theta(x)_y\big)}{\sum_{j=1}^{k} \exp\big(h_\theta(x)_j\big)}\right) =$

    $\log\big(\sum_{j=1}^{k} \exp\big(h_\theta(x)_j\big)\big) - h_\theta(x)_y$

  - Minimizing $\text{Loss}(h_\theta(x), y)$ amounts to maximizing the logit $\big(h_\theta(x)\big)_y$ corresponding to the correct label y

Cross-Entropy Loss Example

|	Cat	Dog	Fox	Cow	Red Panda	Bear	Dolphin
True distribution:	0%	0%	0%	0%	100%	0%	0%
Predicted distribution:	2%	30%	45%	0%	25%	5%	0%

Classifier

Cross-Entropy Loss:

$$H(\mathbf{p}, \mathbf{q}) = -\Sigma_i\, p_i \log(q_i)$$
$$= -\log(0.25) = 1.386$$

$\log_2(x) = \log(x) / \log(2)$

Cross-Entropy Loss Example

- Consider a NN for 3-class classification. Fig shows the last linear layer and the SoftMax layer

- The last linear layer computes the vector of logits $h_\theta(x) = Wx_i + b = [-2.85 \quad .86 \quad .28]^T$ (x is the input image to the NN, x_i is the intermediate input to the last layer)

- The SoftMax layer computes the vector of predicted probabilities $[.016 \quad .631 \quad .353]^T$ for labels $[1 \quad 2 \quad 3]^T$, and the loss $-\log .353$, assuming correct label $y_i = 3$

Last Linear Layer   SoftMax Layer
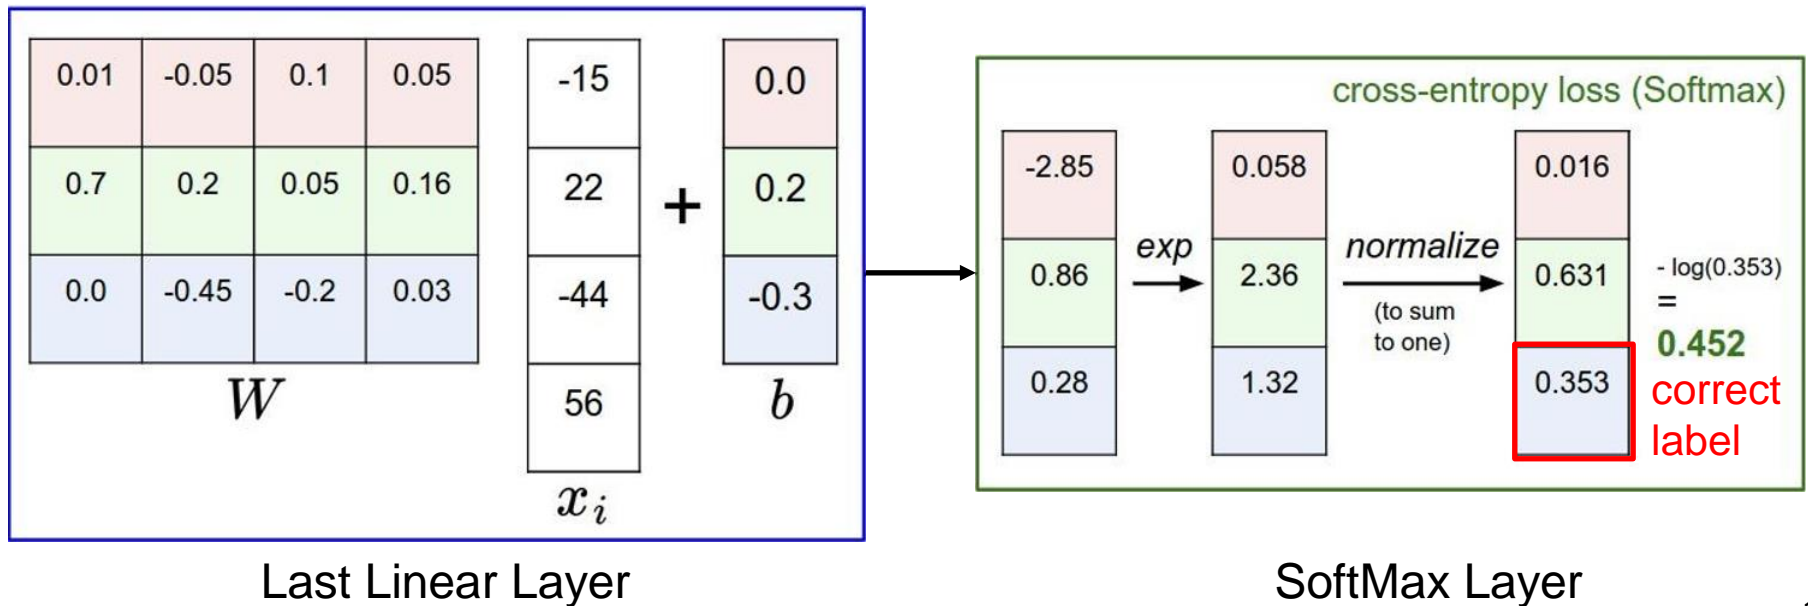
Example CV Task: Multi-Class Image Classification

- Two stages: feature extraction from input, and classification based on extracted features

- Classifier returns output as a list of probabilities with size equal to the number of classes, but it may also return the top-1 or top-5 results with highest probability ranking

$P(0) = .01$
$P(1) = .01$
$P(2) = .01$
$P(3) = .02$
$P(4) = .03$ sum
$P(5) = .01$ to 1.0
$P(6) = .02$
$P(7) = .02$
$P(8) = .85$
$P(9) = .02$

Feature Extraction

Classifier

CNN
(or classic CV algo)

SoftMax
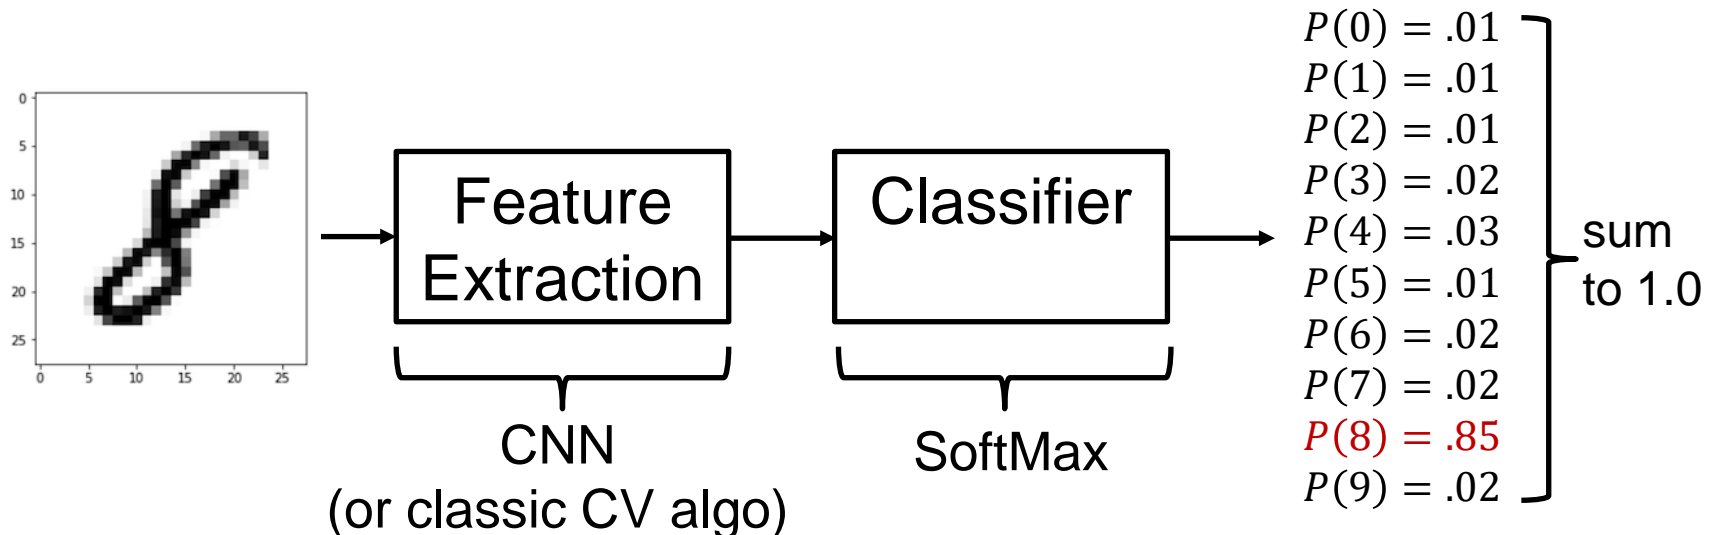
Binary Classification Metrics

- The relevant class is considered "positive" in a binary classifier
- e.g., for a medical test that aims to diagnose people with a certain disease. "Positive" denotes sick (has disease), and "negative" denotes healthy (no disease)
  - TP: a sick person is diagnosed as sick
  - TN: a healthy person is diagnosed as healthy
  - FP: a healthy person is misdiagnosed as sick
  - FN: a sick person is misdiagnosed as healthy

relevant elements

false negatives | true negatives

true positives | false positives

selected elements
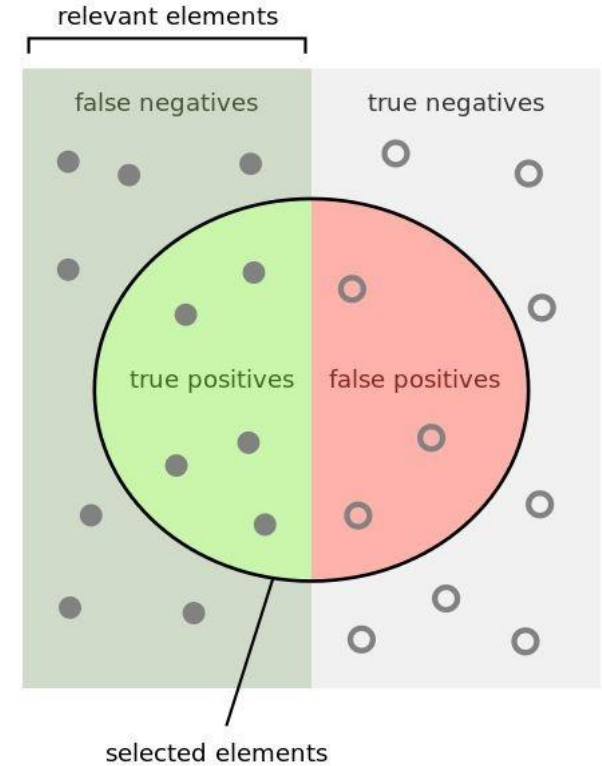
How many selected items are relevant?

How many relevant items are selected?

$Precision = \dfrac{}{}$

$Recall = \dfrac{}{}$

Example Confusion Matrix 1

- Precision $= \frac{TP}{TP+FP} = \frac{1}{1+7} = .125$
  - When the classifier predicts positive, it is correct 12.5% of the time
- Recall (TPR) $= \frac{TP}{TP+FN} = \frac{1}{1+2} \approx .333$
  - Among all the positive cases, the classier correctly classifies 33.3% of them as positive

- $F1 = 2 * \frac{\text{Precision}*\text{Recall}}{(\text{Precision}+\text{Recall})} = 2 * \frac{.333*.125}{.333+.125} = .182$

- False Positive Rate (FPR) $= \frac{FP}{FP+TN} = \frac{7}{7+90} \approx .072$
  - Among all the negative cases, the classier misclassifies 7.2% of them as positive
- Accuracy $= \frac{TP+TN}{TP+TN+FP+FN} = \frac{1+90}{1+90+7+2} = .91$
  - The classier makes the correct prediction 91% percent of the time
- Positive correlation between TPR vs. FPR; Possibly non-monotonic correlation between precision vs. recall

		Ground Truth	
		Positive	Negative
Predicted	Neg	False Negative (FN)=2	True Negative (TN)=90
	Pos	True Positive (TP)=1	False Positive (FP)=7

https://medium.com/swlh/recall-precision-f1-roc-auc-and-everything-542aedf322b9
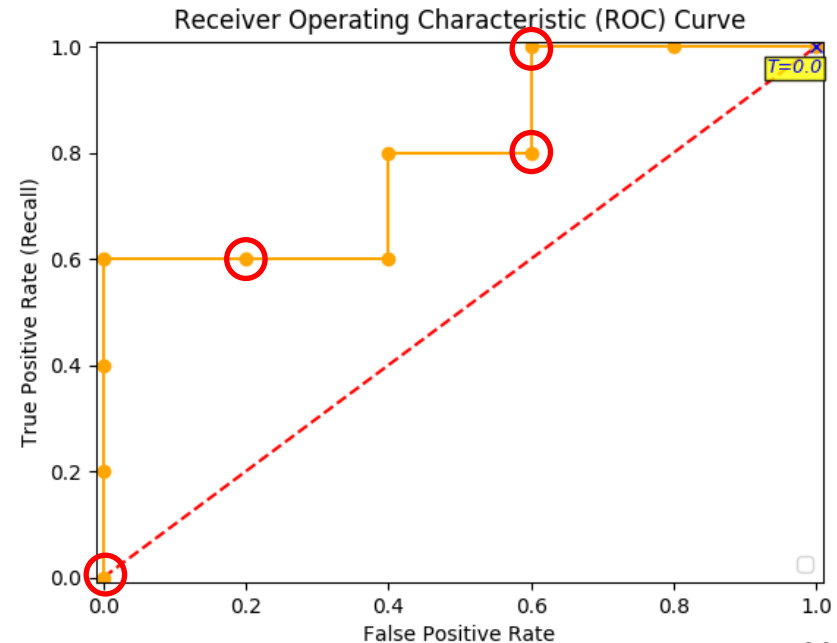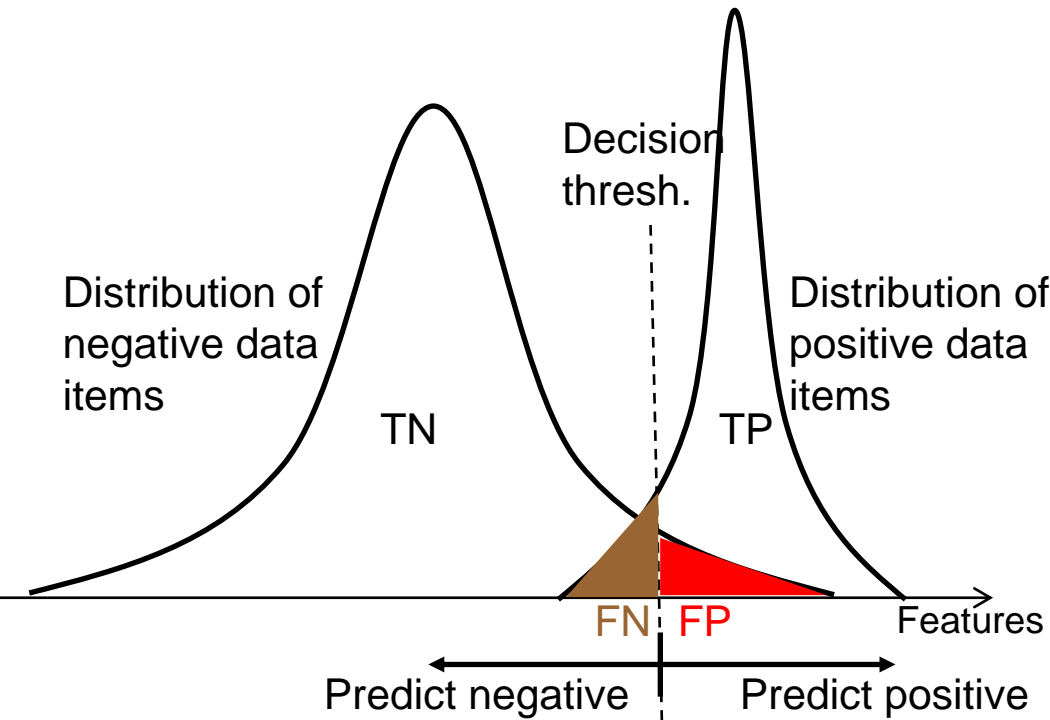
21

Example Confusion Matrix 2

- Precision = $\frac{TP}{TP+FP} = \frac{0}{0+0}$ (ill-defined)
  - When the classifier predicts positive, it is correct ?% of the time (since it never predicts positive, the question is ill-defined)
- Recall (TPR) = $\frac{TP}{TP+FN} = \frac{0}{0+3} = 0$
  - Among all the positive cases, the classier correctly classifies 0% of them as positive
- False Positive Rate (FPR) = $\frac{FP}{FP+TN} = \frac{0}{0+97} = 0$
  - Among all the negative cases, the classier misclassifies 0% of them as positive
- Accuracy = $\frac{TP+TN}{TP+TN+FP+FN} = \frac{0+97}{0+97+0+3} = .97$
  - The classier makes the correct prediction 97% percent of the time
- A medical test that never makes any positive diagnoses is very accurate for a rare disease (diagnose everyone to be healthy), but not very useful

		Ground Truth	
		Positive	Negative
Predicted	Neg	False Negative (FN)=3	True Negative (TN)=97
	Pos	True Positive (TP)=0	False Positive (FP)=0

ROC and AUC

- Binary classification is typically based on a decision threshold parameter. Moving the decision threshold will cause FPR and TPR to move in the same direction
  - e.g., a medical test that sets a lower threshold for positive diagnosis will have both higher FPR and higher TPR, and vice versa
- Receiver Operating Characteristic (ROC) Curve plots FPR (x-axis) vs. TPR (y-axis); Area Under the Curve (AUC) is the area under ROC ($.5 \leq ROC \leq 1$, since $FPR \leq TPR$)
  - Fig shows an example with 4 points (FPR, TPR) highlighted: (0,0), (.2, .6), (.6, .8), (.6,1.0)
  - The ideal ROC curve: $FPR \equiv 0, TPR \equiv 1, AUC = 1$, with $FP = FN = 0$,
  - The worst ROC curve; $FPR \equiv TPR, AUC = .5$ (dotted line)

Decision thresh.

Distribution of negative data items

Distribution of positive data items

TN

TP

FN FP Features

Predict negative Predict positive

23

Confusion Matrix for Multi-Class Classification

- Binary classification is a special case of multi-class classification:

$$Accuracy = \frac{\sum_{i=1}^{3} x(i,i)}{\sum_{i=1}^{3} \sum_{j=1}^{3} x(i,j)}$$

Binary ⟹

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

		Ground Truth		
		Cls1	Cls2	Cls3
Pred.	Cls3			
	Cls2			
	Cls1			

|		Ground Truth	
		Pos	Neg
Pred.	Pos	FN	TN
	Neg	TP	FP

K-Fold Cross-Validation

- Divide data into train data and test data
- Since we cannot peek at the test data during training time, we use part of the train data for Cross-Validation:
- e.g., Divide training data into K=5 parts (folds). Use each fold as validation data, and the other 4 folds as training data. Cycle through the choice of which fold used for validation and average results.
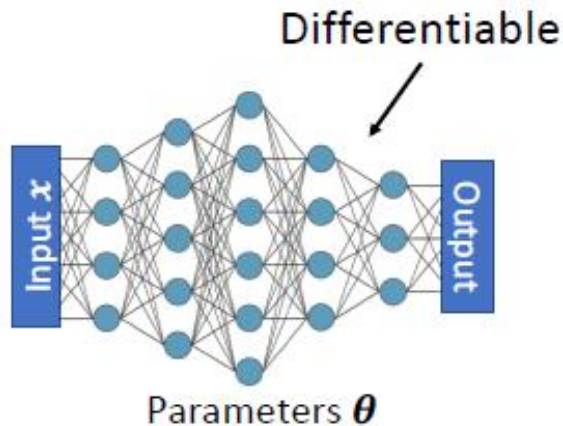
Training Neural Networks

Local Gradient at One Neuron

x

activations

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial x}$$

- "local gradient"

$$\frac{\partial z}{\partial x}$$

z

• f

$$\frac{\partial L}{\partial z}$$

$$\frac{\partial z}{\partial y}$$

y

$$\frac{\partial L}{\partial y} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial y}$$
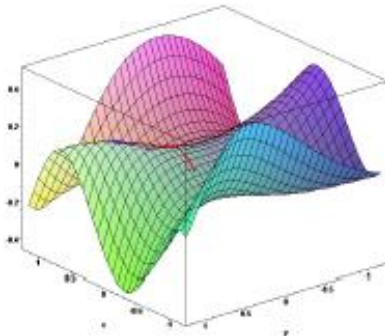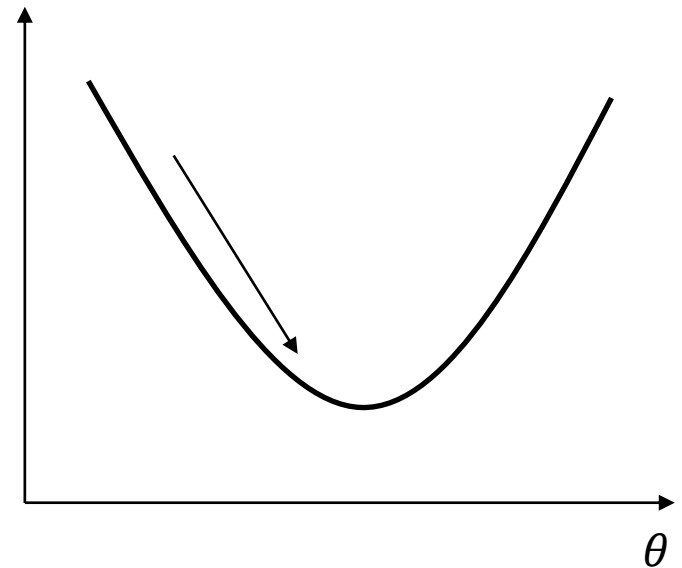
- gradients

Gradient Descent

- Gradient descent $\theta \leftarrow \theta - \alpha \nabla_\theta \text{Loss}(x, y; \theta)$
- Loss surface of a DNN is highly non-convex; can only hope to find "reasonably good" local minima

Differentiable

Input x

Output

Parameters $\boldsymbol{\theta}$

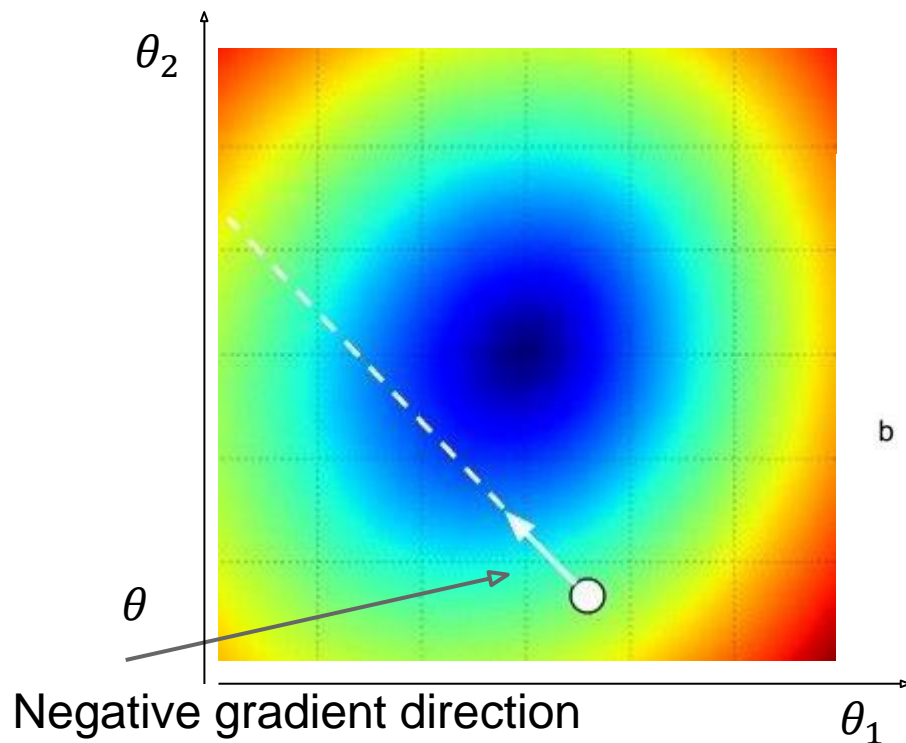Can use gradient descent method to find good θ

$$\mathbb{E}_{(x,y) \sim D} \text{Loss}(x, y; \theta)$$

θ

Gradient Descent Algorithms

- Steepest descent may result in in efficient zig-zag path
- More advanced GD methods exploit momentum, e.g., Nesterov, AdaGrad, RMSProp, Adam…

Negative gradient direction

Mini-batch Stochastic Gradient Descent

- Only use a small portion (a mini-batch) of the training data to compute the gradient
- Common mini-batch sizes are 32/64/128 examples
- Loop:
  - Sample a mini-batch of data
  - Forward prop it through the graph, get loss
  - Backprop to calculate the gradients
  - Update the parameters using gradient descent

Batch Normalization

- For each mini-batch:
  - 1. Compute the empirical mean and variance independently for each dimension $i = 1, \ldots m$
  - 2. Normalize to a unit Gaussian with 0 mean and unit variance
- BN layers inserted before nonlinear activation function, and it keeps x's average value around 0 for maximum gradient during learning
- Scale and shift params γ, β gives more flexibility during training
- Benefits:
  - Improves gradient flow through the network; Allows higher learning rates; Reduces the strong dependence on initialization; Acts as a form of regularization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\ldots m}\}$;
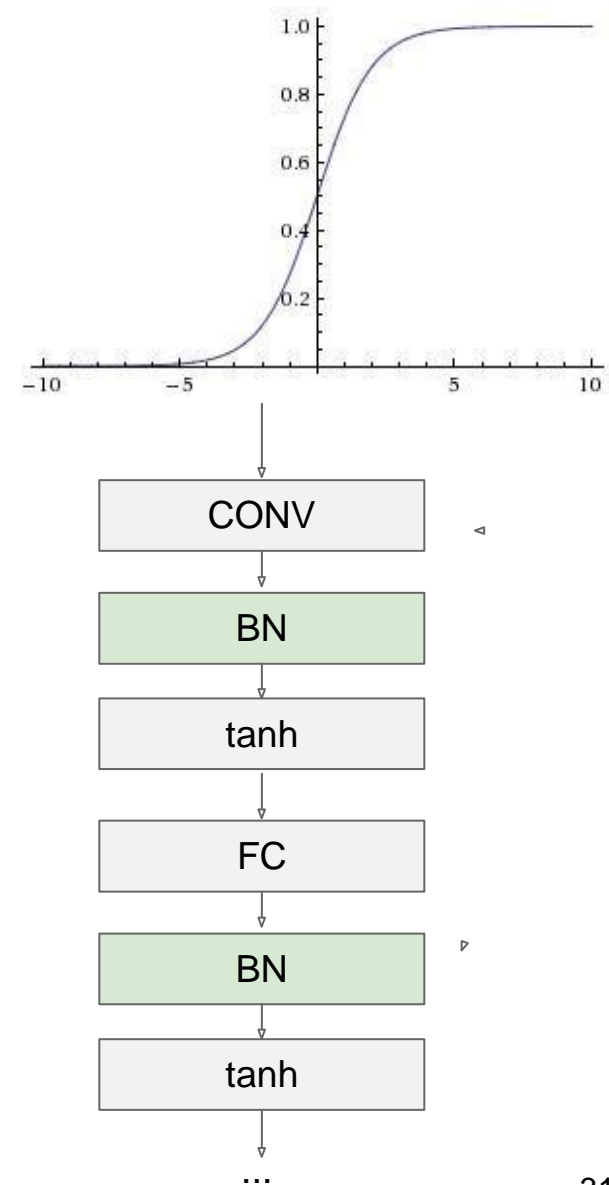Parameters to be learned: γ, β

Output: $\{y_i = \mathrm{BN}_{\gamma,\beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^{m} x_i \qquad \text{// mini-batch mean}$$
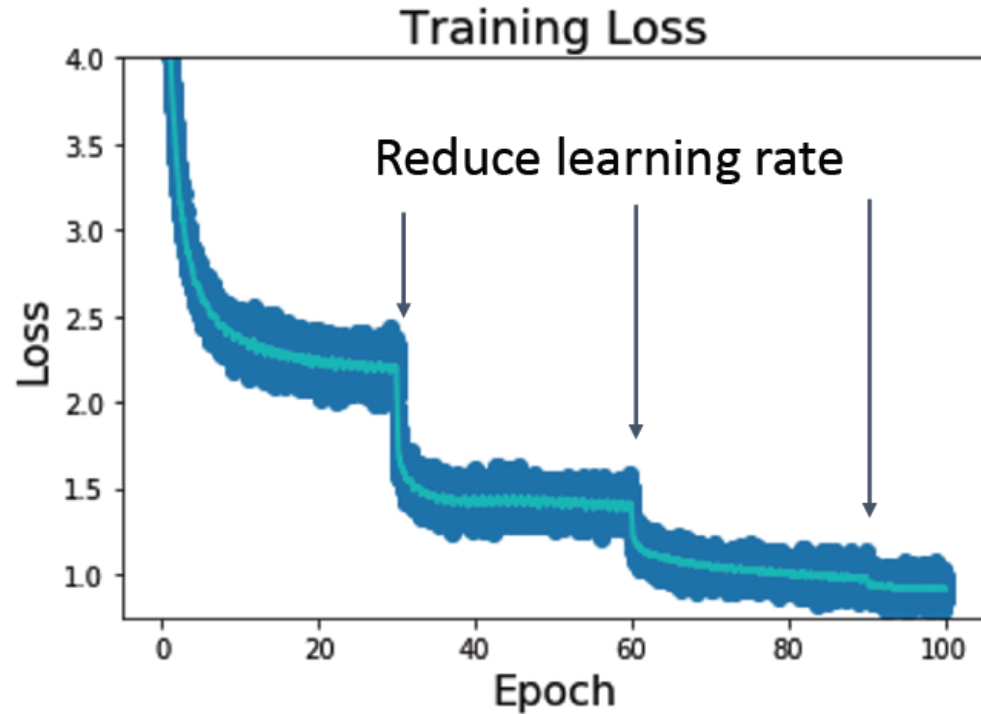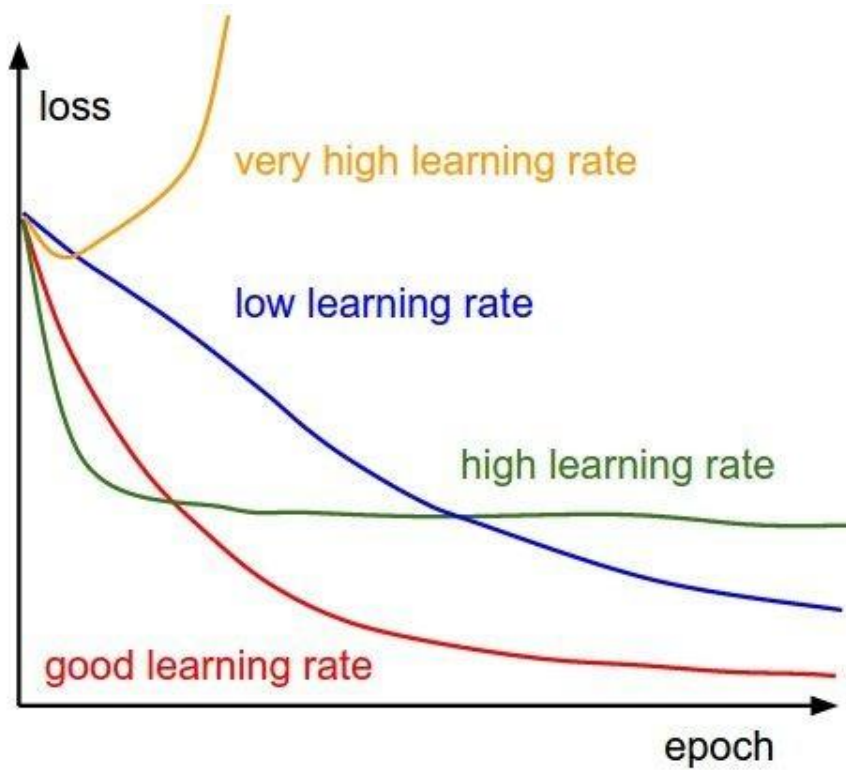
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^{m} (x_i - \mu_{\mathcal{B}})^2 \qquad \text{// mini-batch variance}$$

$$\widehat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \qquad \text{// normalize}$$

$$y_i \leftarrow \gamma \widehat{x}_i + \beta \equiv \mathrm{BN}_{\gamma,\beta}(x_i) \qquad \text{// scale and shift}$$

CONV

BN

tanh

FC

BN

tanh

...

31

# Learning Rate Schedule during Training

Hyperparameter Optimization

- Example hyperparams
  - Network architecture
  - Learning rate, its decay schedule, update type
  - Regularization (L2/Dropout strength)
- Grid search vs. random search
  - Random search can use the computing budget more effectively
  - With 9 evaluations, random search explored 9 different values for the important parameter; grid search only explored 3.

Grid Layout

Random Layout

Unimportant parameter

Important parameter

Classification Accuracy

- Big gap between training accuracy and validation accuracy may imply overfitting => decrease model capacity?
- No gap may imply underfitting => increase model capacity?
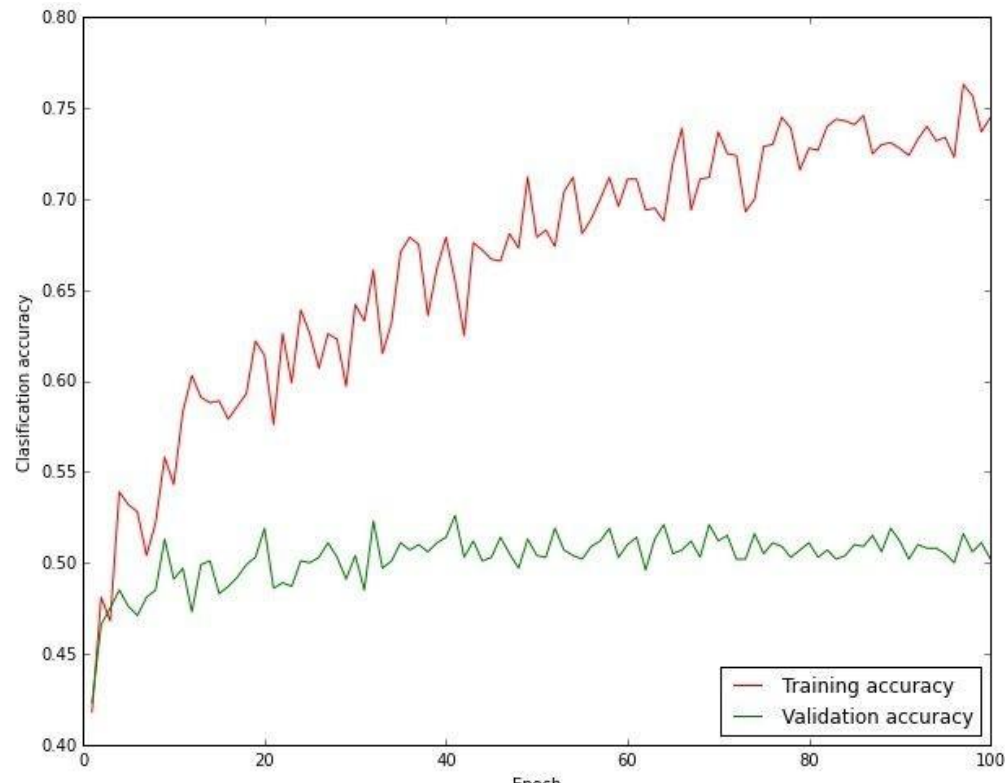
Data Augmentation for Enlarging Training Dataset

- Mirroring, random cropping, color shifting, rotation, shearing, local warping…

Mirroring

Random Cropping

Color Shifting