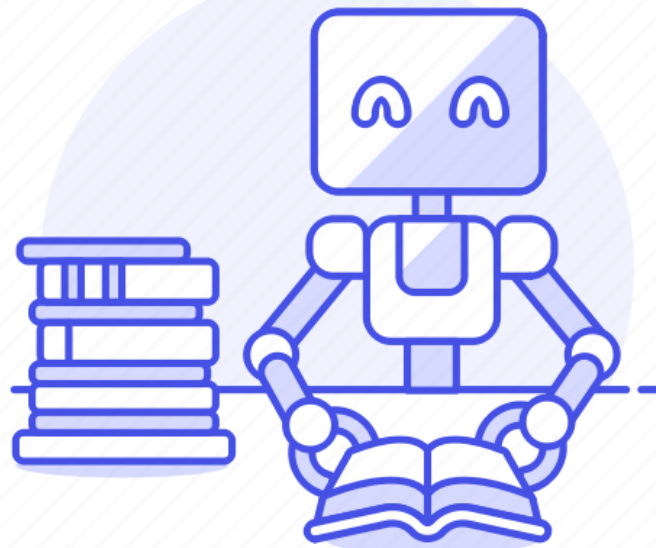


L3 Introduction to Machine Learning

Z. Gu 2021




ML Taxonomy

- Supervised Learning:
 - The system is presented with example inputs and their desired outputs, given by a “teacher”, and the goal is to learn a general rule that maps inputs to outputs.
 - Classification (cat or dog?)
 - Regression (housing price next year?)
- Unsupervised Learning:
 - No labels are given to the learning algorithm, leaving it on its own to find structure in its input. Unsupervised learning can be a goal in itself (discovering hidden patterns in data) or a means towards an end (feature learning).
 - Parametric UL (e.g., Gaussian Mixture Models)
 - Non-parametric UL


ML Taxonomy

- Reinforcement Learning:
 - An agent interacts with a dynamic environment in which it must perform a certain goal. The agent is provided feedback in terms of rewards and it tries to learn an optimal policy that maximizes its cumulative rewards.
 - Algorithms: Model-based; Model-free (Value-based, Policy-based)
 - Applications: Game playing (AlphaGo); Robotics; AD...



Training vs. Inference

- Training: millions of iterations of forward pass + back propagation to adjust model params (e.g., NN weights); requires large CPU/GPU clusters and days/weeks of training time
- Inference (also called prediction): a single forward pass; can be run on edge devices



Supervised Learning for Image Classification

- Collect a dataset of images and labels
- Use Machine Learning to train an image classifier
- Evaluate the classifier on test images


```
def train(train_images, train_labels):  
    # build a model for images -> labels...  
    return model  
  
def predict(model, test_images):  
    # predict test_labels using the model...  
    return test_labels
```

Example training set



Image Classification

- Input encoding:
 - A color image is represented as a 3D tensor, with integers between [0, 255] denoting pixel intensities, e.g. 300x100x3 (3 color channels RGB)
 - A greyscale image is represented as a 2D tensor, e.g. 300x100x1
- Two stages: feature extraction from input, and classification based on extracted features
- Classifier returns output as a list of probabilities with size equal to the number of classes, but it may also return the top-1 or top-5 results with highest probability ranking



Feature
Extraction

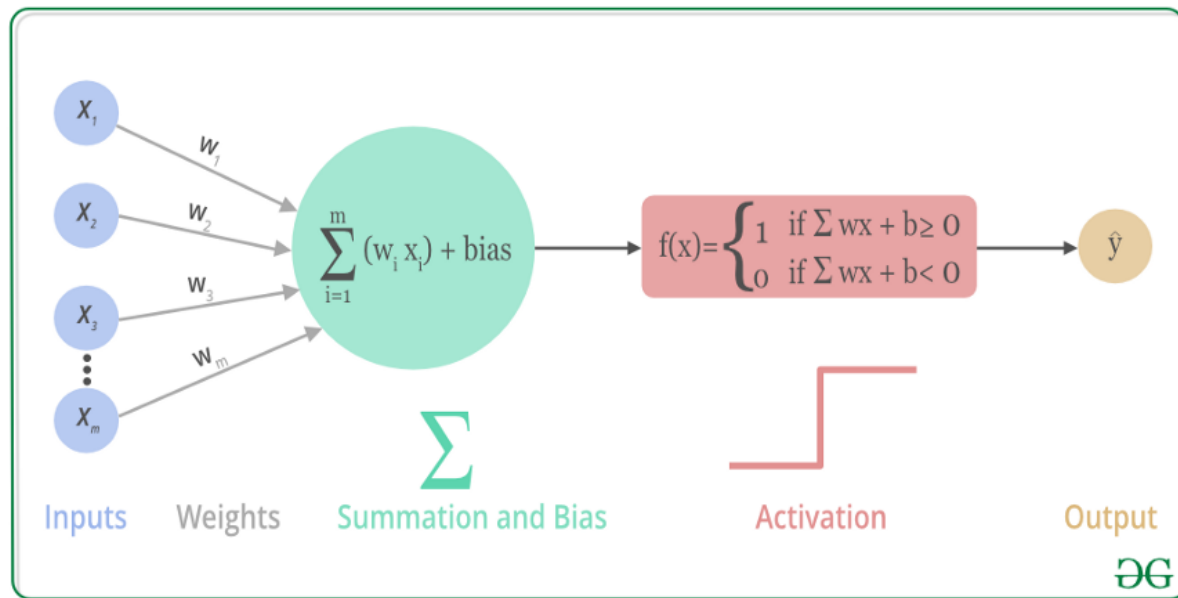
Classifier

$P(0) = .01$
 $P(1) = .01$
 $P(2) = .01$
 $P(3) = .02$
 $P(4) = .03$
 $P(5) = .01$
 $P(6) = .02$
 $P(7) = .02$
 $P(8) = .85$
 $P(9) = .02$

} sum
to 1.0


A Neuron and its Activation Function

- A neuron may perform:
 - Linear Regression if $y = a = wx + b$ (activation function $f =$ identity function)
 - Logistic Regression if $y = \sigma(a) = \sigma(wx + b)$ (activation function $f =$ sigmoid function, used for binary classification)
 - Perceptron classification $y = \sigma(a) = \text{step}(wx + b)$ (activation function $f =$ step function, shown below)
- The activation function is a nonlinear monotonic function that acts like a “gate”: the output is larger for larger input activation
 - Many variants: step, sigmoid, tanh, ReLU, leaky-ReLU, PReLU, SoftPlus, Swish...




Linear Regression for Regression

- Function approximation $y = wx + b$, with learnable parameters $\theta = \{w, b\}$, where x, y, b are vectors, and w is a weight matrix
 - e.g., we want to predict price of a house based on its feature vector $\mathbf{x} = [x_1 \ x_2 \ x_3]^T$, where x_1 is area in square meters (sqm), x_2 is location ranking (loc), x_3 is year of construction (yoc)
 - Predicted price $y = wx + b = w_1x_1 + w_2x_2 + w_3x_3 + b$
 - Fig shows an example for scalar x and y







Logistic Regression for Binary Classification

- Consider a binary classification problem: an input image x may be classified as a dog with probability $P(y = dog|x)$, a cat with probability $P(y = cat|x)$, with $P(y = dog|x) + P(y = cat|x) = 1.0$
- Logistic Regression: use sigmoid function $\sigma(a) = \frac{1}{1+e^{-a}}$ to map from activation (also called the logit) to output probabilities



Common Activation Functions used in DL




Deep Neural Networks

- We can stack many hidden layers to form a DNN if we have enough data and computing power to train it
- The high model capacity of DNN comes from non-linear mappings: hidden units must be followed by a non-linear activation function
 - Without non-linear activation functions, a DNN with many layers can be collapsed into an equivalent single-layer NN

Fully-Connected NNs

- Number of params to learn at i-th layer is $(N_{i-1} + 1) * N_i$, where N_i is the number of neurons at i-th layer.
Can grow very large
 - (We will discuss CNNs in the next lecture with much fewer params)




Slide Credit: Hugo Laroché NN course


A 3-layer NN

Example: Two-Layer Fully-Connected NN for Solving XOR

- The NN consists of one input, one hidden, and one output layer, with sigmoid activations




x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



Setting # Layers and Their Sizes

- An example illustrating adding more hidden neurons increases model capacity and reduces training error
- But too many layers and neurons may lead to overfitting



NN for Multi-Class Classification

- Consider a NN defining the model, $h_\theta: \mathcal{X} \rightarrow \mathbb{R}^k$, as the mapping from input x to output $h_\theta(x)$, a k -dim vector of logits, where k is the number of classes
 - θ is the set of params (weights and biases)
 - y is the correct label for input x
 - Note that h_θ does not include the last SoftMax layer
- e.g., a 3-layer NN consisting of 2 layers with ReLU activation functions and a last linear layer is
 - $h_\theta(x) = W_3 \max(0, W_2 \max(0, W_1 x + b_1) + b_2) + b_3$

Cross-Entropy Loss for Multi-Class Classification

- The SoftMax operator $\sigma: \mathbb{R}^k \rightarrow \mathbb{R}^k$ computes a vector of predicted probabilities $\sigma(z): \mathbb{R}^k$ from a vector of logits $z: \mathbb{R}^k$, where k is the number of classes:

- $\sigma(z)_i = \frac{\exp(z_i)}{\sum_{j=1}^k \exp(z_j)}$

- The loss function is defined as the negative log likelihood of the predicted probability corresponding to the correct label y :

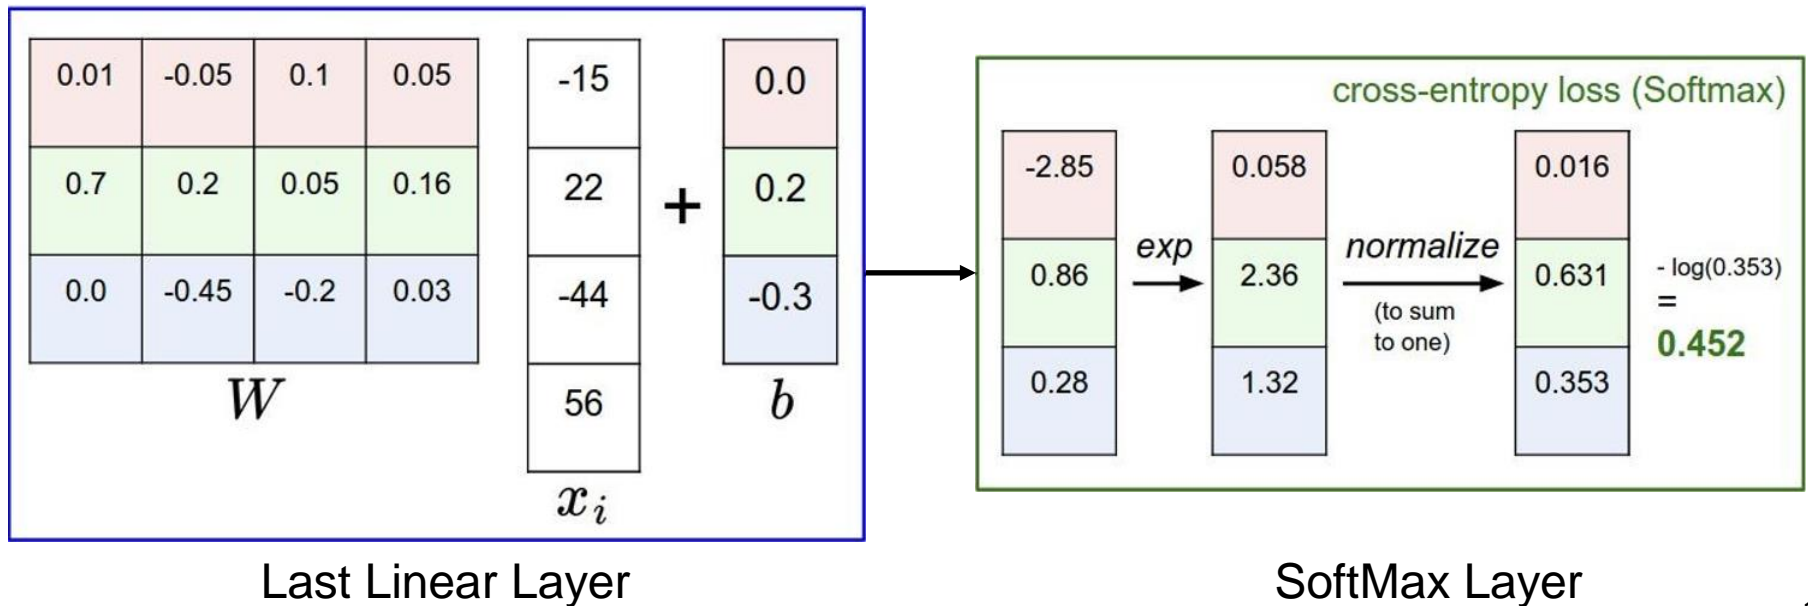
- $\text{Loss}(h_\theta(x), y) = -\log \sigma(h_\theta(x))_y = -$

- $\log \left(\frac{\exp(h_\theta(x)_y)}{\sum_{j=1}^k \exp(h_\theta(x)_j)} \right) = \log(\sum_{j=1}^k \exp(h_\theta(x)_j)) - h_\theta(x)_y$

- Minimizing $\text{Loss}(h_\theta(x), y)$ amounts to maximizing the logit $(h_\theta(x))_y$ corresponding to the correct label y

Cross-Entropy Loss Example

- Consider a NN for 3-class classification. Fig shows the last linear layer and the SoftMax layer
- The last linear layer computes the vector of logits $h_{\theta}(x) = Wx_i + b = [-2.85 \quad .86 \quad .28]^T$ (x is the input image to the NN, x_i is the intermediate input to the last layer)
- The SoftMax layer computes the vector of predicted probabilities $[.016 \quad .631 \quad .353]^T$ for labels $[1 \quad 2 \quad 3]^T$, and the loss $-\log .353$, assuming correct label $y_i = 3$



Binary Classification Metrics


- The relevant class is considered “positive” in a binary classifier
- e.g., for a medical test that aims to diagnose people with a certain disease. “Positive” denotes sick (has disease), and “negative” denotes healthy (no disease)
 - TP: a sick person is diagnosed as sick
 - TN: a healthy person is diagnosed as healthy
 - FP: a healthy person is misdiagnosed as sick
 - FN: a sick person is misdiagnosed as healthy

		Prediction	
		Positives	Negatives
Ground Truth	Positives	True positive (TP)	False negative (FN)
	Negatives	False positives (FP)	


$$\text{Recall} = \frac{TP}{TP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$


$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$



How many selected items are relevant?




How many relevant items are selected?



ROC Curve


- Binary classification results typically depend on a parameter (eg. decision threshold) whose value impacts precision and recall.
- Receiver Operating Curve (ROC curve):
 - x-axis: False Positive Rate $FPR = \frac{FP}{FP+TN}$ = recall
 - y-axis: True Positive Rate $TPR = \frac{TP}{TP+FN}$
 - The ideal ROC curve: $FPR \equiv 0, TPR \equiv 1$, with $FP = FN = 0$
- Moving the decision threshold will cause FPR and TPR to move in the same direction
 - e.g., a medical test that sets a high threshold for positive diagnosis will have both low FPR and low TPR, and vice versa



Confusion Matrix for Multi-Class Classification

- Binary classification as a special case of multi-class classification:

$$Accuracy = \frac{\sum_{i=1}^3 x(i, i)}{\sum_{i=1}^3 \sum_{j=1}^3 x(i, j)}$$



$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

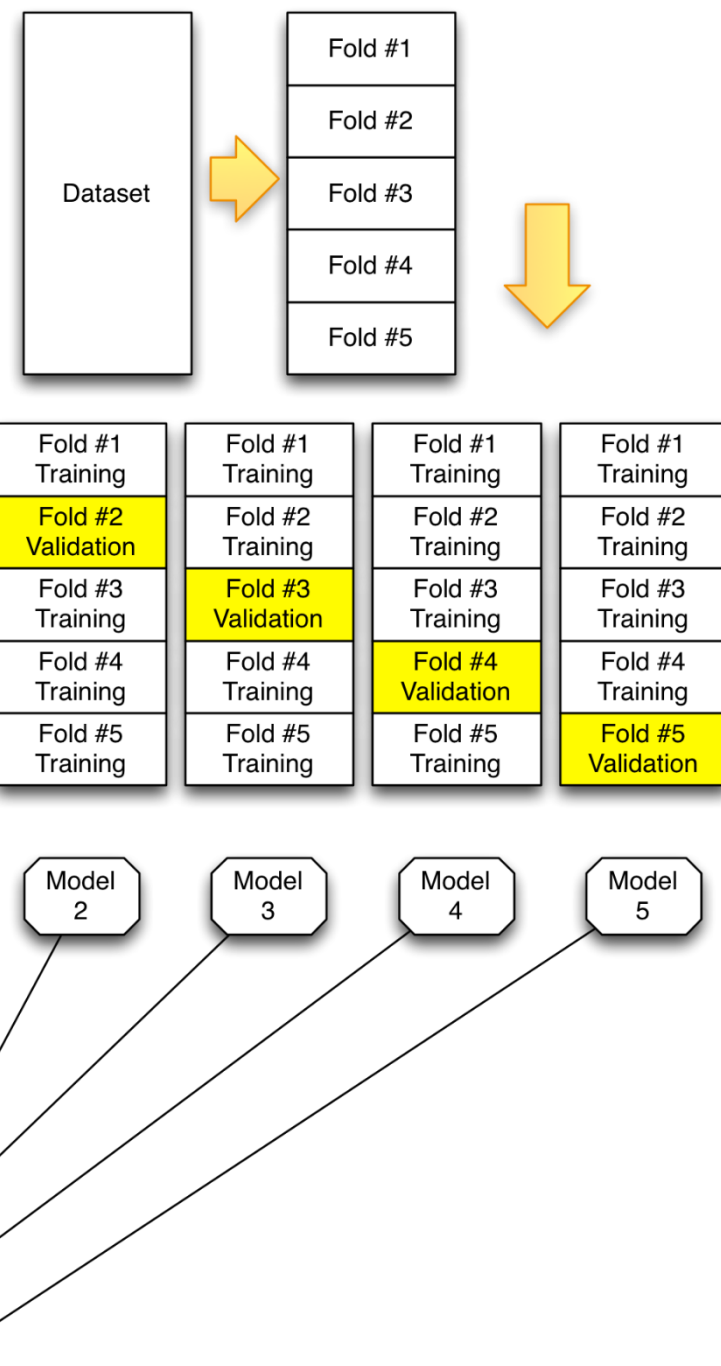
		Prediction		
		Class 1	Class 2	Class 3
Ground Truth	Class 1	x(1,1)	x(1,2)	x(1,3)
	Class 2	x(2,1)	x(2,2)	x(2,3)
	Class 3	x(3,1)	x(3,2)	x(3,3)

		Prediction	
		Positives	negative
Ground Truth	Positives	True positive (TP)	False negative (FN)
	Negative	False positives (FP)	True negative (TN)

Training Neural Networks


K-Fold Cross-Validation

- Divide data into train data and test data
- Since we cannot peek at the test data during training time, we use part of the train data for Cross-Validation:
- e.g., Divide training data into $K=5$ parts (folds). Use each fold as validation data, and the other 4 folds as training data. Cycle through the choice of which fold used for validation and average results.





Gradient Descent

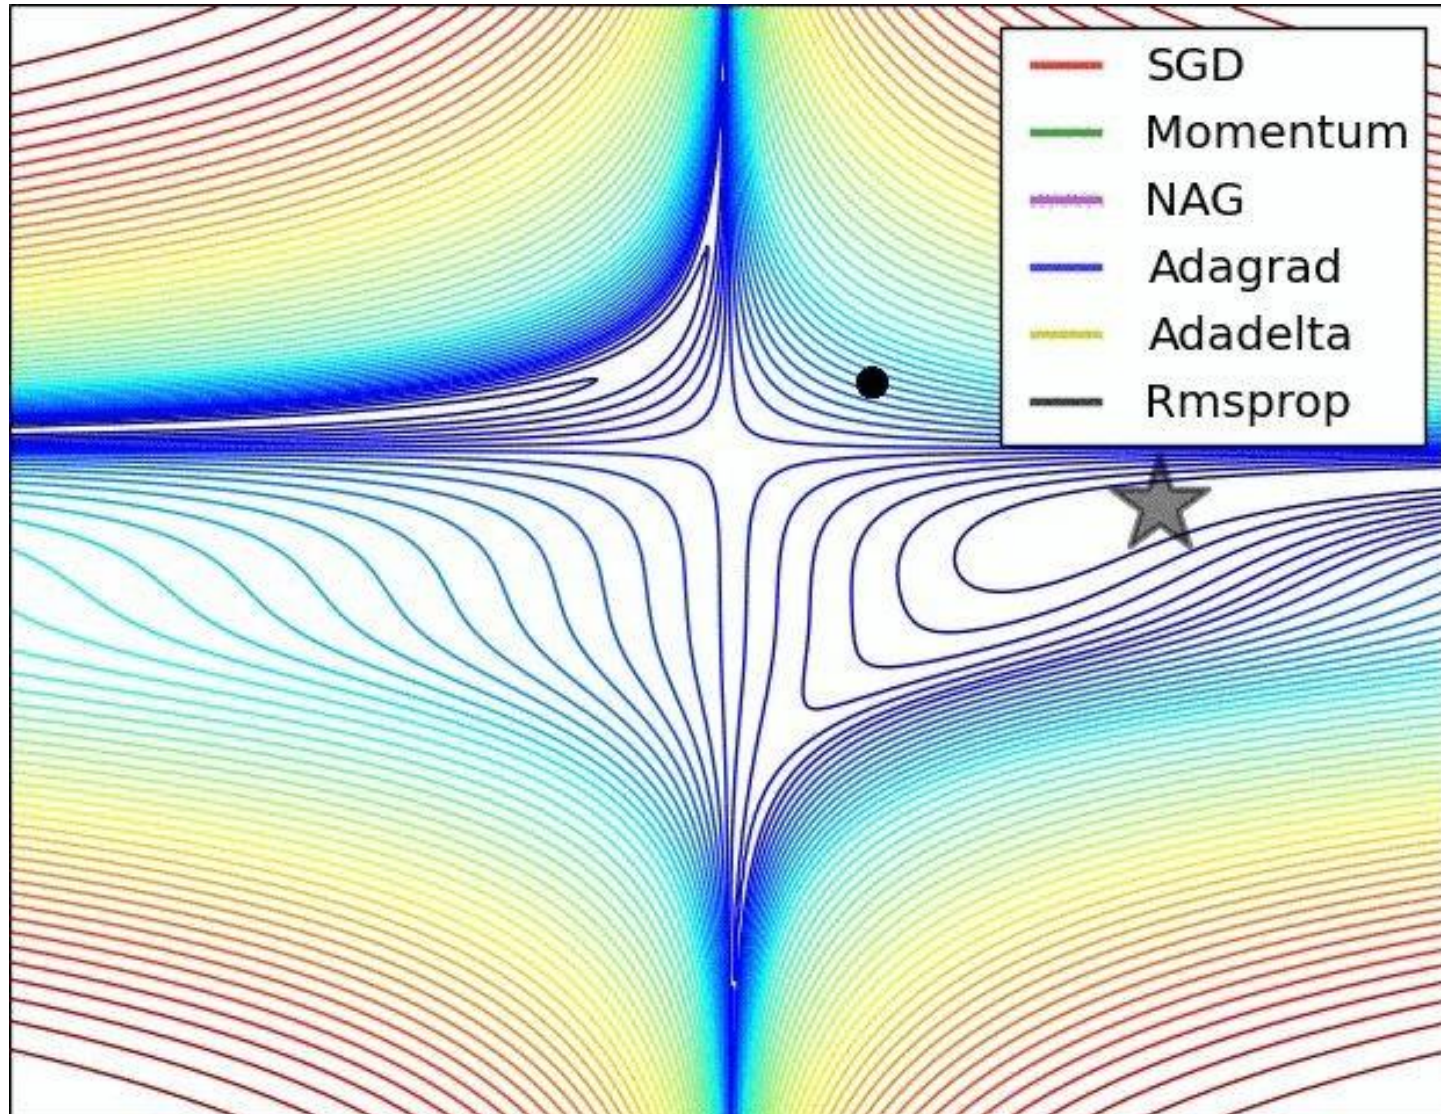
- Gradient descent $\theta \leftarrow \theta - \alpha \nabla_{\theta} \text{Loss}(x, y; \theta)$



Can use gradient descent method to find good θ



Different GD Update Formulas



Mini-batch Stochastic Gradient Descent

- Only use a small portion (a mini-batch) of the training data to compute the gradient
- Common mini-batch sizes are 32/64/128 examples
- Loop:
 - Sample a mini-batch of data
 - Forward prop it through the graph, get loss
 - Backprop to calculate the gradients
 - Update the parameters using gradient descent

Batch Normalization

- For each mini-batch:
 - 1. Compute the empirical mean and variance independently for each dimension $i = 1, \dots, m$
 - 2. Normalize to a unit Gaussian with 0 mean and unit variance
- BN layers inserted before nonlinear activation function, and it keeps x 's average value around 0 for maximum gradient during learning
- Scale and shift params γ, β gives more flexibility during training
- Benefits:
 - Improves gradient flow through the network; Allows higher learning rates; Reduces the strong dependence on initialization; Acts as a form of regularization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1, \dots, x_m\}$;

Parameters to be learned: γ, β


Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$


$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$





Hyperparameter Optimization

- Example hyperparams
 - Network architecture
 - Learning rate, its decay schedule, update type
 - Regularization (L2/Dropout strength)
- Grid search vs. random search



The Loss Curve during Training



Classification Accuracy

- Big gap between training accuracy and validation accuracy may imply overfitting => decrease model capacity?
- No gap may imply underfitting => increase model capacity?

