

# Lecture 9

## Balanced Search Trees

Jianchen Shan  
Department of Computer Science  
Hofstra University

# Lecture Goals

- Develop **balanced search trees** with guaranteed logarithmic performance for search and insert (and many other operations).
- We begin with **2-3 trees**, which are easy to analyze but hard to implement.
- Next, we consider **red-black** binary search trees, which we view as a novel way to implement 2-3 trees as binary search trees.
- Finally, we introduce **B-trees**, a generalization of 2-3 trees that are widely used to implement file systems.

# 2-3 Search Trees

The 2-3 tree is a way to generalize BSTs to provide the flexibility that we need to guarantee fast performance

Allow 1 or 2 keys per node


- 2-node: one key, two children.
- 3-node: two keys, three children.

Perfect balance: Every path from root to null link has same length.

Symmetric order: Inorder traversal yields keys in ascending order.

← how to maintain?

So the search is a generalization of the search in BSTs



## Search

- Compare search key against keys in node.
- Find interval containing search key.
- Follow associated link (recursively).

## Search for H

H is less than M (go left)

H is between E and J (go middle)

Search hit

## Search for B

B is less than M (go left)

B is less than E (go left)

B is between A and C (go middle)

Link is null

Search miss

# Insert in 2-3 Search Trees

Insertion into a 2-node at bottom.

- Search for key, as usual
- Add new key to 2-node to create a 3-node.

Insert K

K is less than M (go left)

K is larger than J (go right)

K is less than L

Search ends here and replace 2-node with 3-node containing K

Insert Z

Z is larger than M (go right)

Z is larger than R (go right)

Z is larger than X


Search ends here

Replace 3-node with temporary 4-node containing Z

Split 4-node into two 2-nodes (pass middle key to parent)

Insertion into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.




# Insert in 2-3 Search Trees (Contd.)

Insertion into a 3-node at bottom.

- Add new key to 3-node to create temporary 4-node.
- Move middle key in 4-node into parent.
- Repeat up the tree, as necessary.
- If you reach the root and it's a 4-node, split it into three 2-nodes.

Insert L


L is between E and R (go middle)




Search ends here

Replace 3-node with temporary 4-node containing L

Split 4-node (move L to parent)



Split 4-node (move L to parent)



Height of tree increases by 1

# 2-3 Search Trees Construction

Insert S Create 2-node in the empty tree

Insert E Convert 2-node into 3-node

Insert A Convert 3-node into 4-node

Split 4-node into two 2-nodes (move E to parent)

Insert R Convert 2-node into 3-node

Insert C Convert 2-node into 3-node

Insert H Convert 3-node into 4-node

Split 4-node into two 2-nodes (move R to parent)

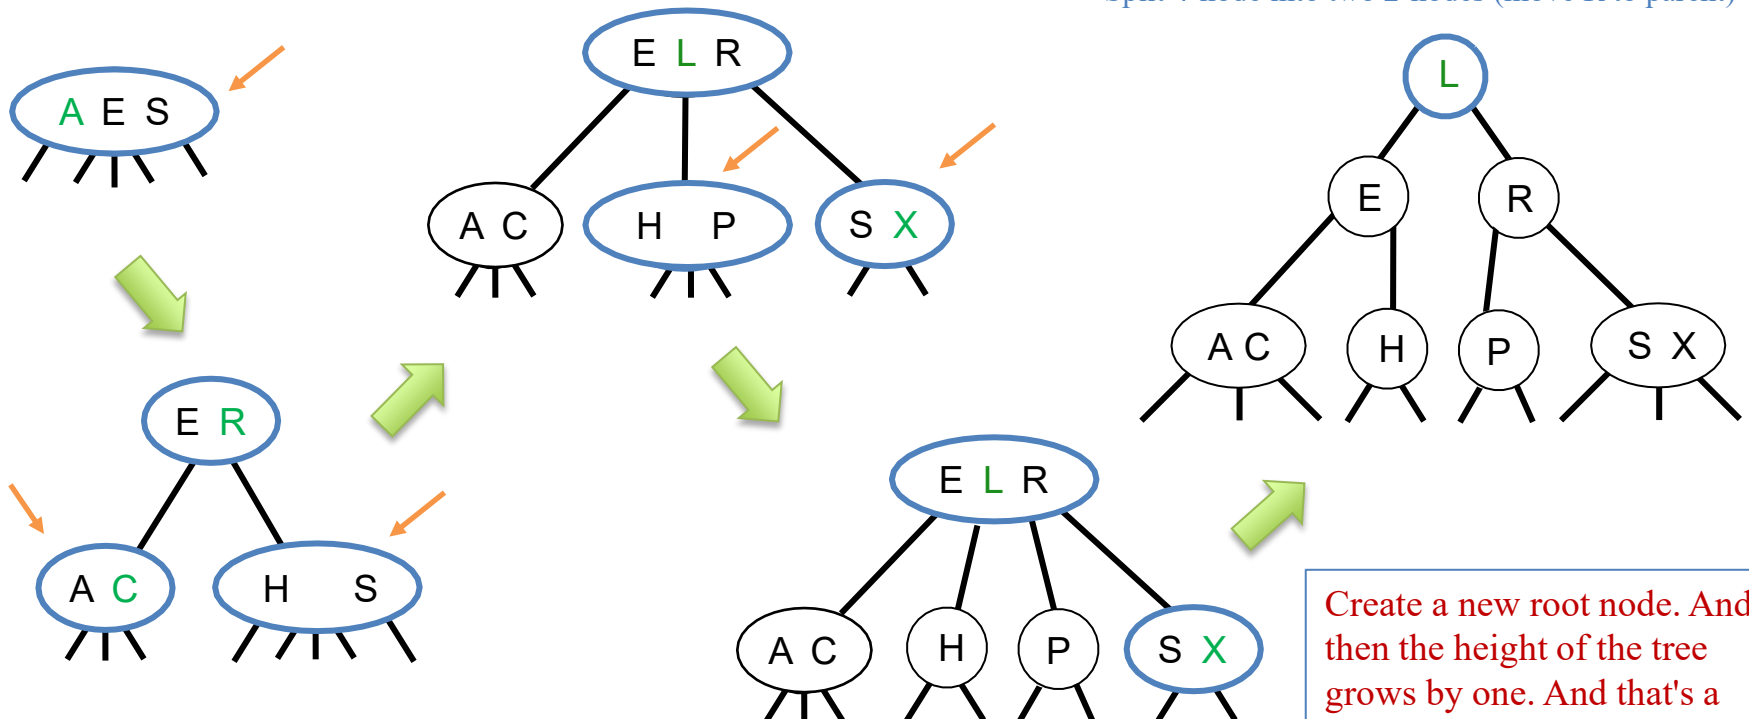
Insert X Convert 2-node into 3-node

Insert P Convert 2-node into 3-node

Insert L Convert 3-node into 4-node

Split 4-node into two 2-nodes (move R to parent)

Split 4-node into two 2-nodes (move R to parent)




Create a new root node. And then the height of the tree grows by one. And that's a legal 2-3 tree, so we stop.

# Local Transformations in a 2-3 Tree

Converting a 2-node to a 3-node

Converting a three to a four, and then splitting and passing a node up



only operations we need to consider to get **balance**.

**local** transformation: constant number of operations.

Only involves changing a constant number of **links**, and is independent of the tree size.

# Global Properties in a 2-3 Tree

**Invariants.** Maintains symmetric order and perfect balance.

**Proof.** Each **transformation** maintains symmetric order and perfect balance.


Splitting a 4-node

root

parent is a 3-node




left




parent is a 2-node


middle




left



right



right




The height doesn't change. If it is perfect balance before, it is perfect balance afterwards



# Performance of 2-3 Tree

**Perfect balance:** Every path from root to null link has same length.



So, that's going to give us, a very easy way to describe a performance.

Operations have costs that's proportional to the path link from root to the bottom.

## Tree height

- **Worst case:**  $\log_2 n$ . [all 2-nodes]
- **Best case:**  $\log_3 n \approx 0.631 \log_2 n$ . [all 3-nodes]
- Between 12 and 20 for a million nodes.
- Between 18 and 30 for a billion nodes.

	Best case	Average case	Worst case
BST	$O(1)$	$O(\log n)$	$O(n)$
2-3 Tree	$O(1)$	$O(\log n)$	$O(\log n)$

2-3 tree model guaranteed **logarithmic** performance for search and insert.

# Lecture Goals

- Develop balanced search trees with guaranteed logarithmic performance for search and insert (and many other operations).
- We begin with 2-3 trees, which are easy to analyze but hard to implement.
- Next, we consider **red-black** binary search trees, which we view as a novel way to implement 2-3 trees as binary search trees.
- Finally, we introduce **B-trees**, a generalization of 2-3 trees that are widely used to implement file systems.

# Red-Black BSTs

## Motivation


Direct implementation of 2-3 trees can be done but it is complicated, because:

- Maintaining multiple node types is cumbersome.
- Need multiple compares to move down tree.
- Need to move back up the tree to split 4-nodes.
- Large number of cases for splitting.

**Red-Black BSTs (Sedgewick 1979):** a simple data structure that allows us to implement 2-3 tree with very little extra code beyond the basic BST code.


**Left-leaning red-black BSTs (Sedgewick 2007) :** the simplest to implement and at least as efficient

- Represent 2–3 tree as a BST.
- Use "internal" left-leaning links as "glue" for 3–nodes.




# Left-Leaning Red-Black BSTs

2-3 tree



red links "glue" nodes within a 3-node

add red link




black links connect 2-nodes and 3-nodes

corresponding red-black BST

merge red link

horizontal red link



Key property. 1-1 correspondence between 2-3 and LLRB.

A BST such that:

- No node has two red links connected to it.
- Every path from root to null link has the same number of black links.
- Red links lean left.

Because the only red links are internal to three nodes.

follows directly from the corresponding property for 2-3 trees

"perfect black balance"


# Search Implementation for Red-Black BSTs

**Observation.** Search is the same as for elementary BST (*ignore color*).

↖ but runs faster because of better balance

```
public Val get(Key key)
{
    Node x = root;
    while (x != null)
    {
        int cmp = key.compareTo(x.key);
        if(cmp < 0) x = x.left;
        else if (cmp > 0) x = x.right;
        else if (cmp == 0) return x.val;
    }
    return null;
}
```

We don't have to change the code at all. Our regular search code doesn't examine the color of a link and so we can just use it exactly as is.



**Remark.** Most other ops (e.g., iteration) are also identical.

We just make sure that we maintain the balance properties during the insertion

# Red-Black BST Representation


Each node is pointed to by precisely one link (*from its parent*)  
→ can encode color of links in nodes.

```
private static final boolean RED    = true;  
private static final boolean BLACK = false;
```

```
private class Node  
{  
    Key key;  
    Value val;  
    Node left, right;  
    boolean color;    // color of parent link  
}
```



```
private boolean isRed(Node x)  
{  
    if (x == null) return false;  
    return x.color == RED;  
}
```

null links are black



# Elementary Red-Black BST Operations

**Left rotation.** Orient a (temporarily) right-leaning red link to lean left.



```
private Node rotateLeft(Node h)
{
    assert isRed(h.right);
    Node x = h.right;
    h.right = x.left;
    x.left = h;
    x.color = h.color;
    h.color = RED;
    return x;  ← new root to connect
               the node above
}
```


local operation that only changes a few links

**Invariants.** Maintains symmetric order and perfect black balance.

# Elementary Red-Black BST Operations (Contd.)


**Right rotation.** Orient a (temporarily) right-leaning red link to lean left.

rotate S right (before)



Their black heights are not changed after rotation.

rotate S right (after)



symmetric code of left rotation

```
private Node rotateRight(Node h)
{
    assert isRed(h.left);
    Node x = h.left;
    h.left = x.right;
    x.right = h;
    x.color = h.color;
    h.color = RED;
    return x;  ← new root to connect
               the node above
}
```


local operation that only changes a few links

**Invariants.** Maintains symmetric order and perfect black balance.




# Elementary Red-Black BST Operations (Contd.)

**Color flip.** Recolor to split a (temporary) 4-node and pass up the center node.



```
private void flipColors(Node h)
{
    assert !isRed(h);
    assert isRed(h.left);
    assert isRed(h.right);
    h.color = RED;
    h.left.color = BLACK;
    h.right.color = BLACK;
}
```





local operation that only changes a few links

**Invariants.** Maintains symmetric order and perfect black balance.

# Insertion in a LLRB Tree: Overview




**Basic strategy.** Maintain 1-1 correspondence with 2-3 trees by applying elementary red-black BST operations.

insert C



corresponding 2-3 tree

*right link red so rotate left*




corresponding 2-3 tree

# Insertion in a LLRB Tree


Warmup 1. Insert into a tree with exactly 1 node.

left




insert a

*search ends at this null link*




- *red link to new node containing a*
- *converts 2-node to 3-node*

right



insert b

*search ends at this null link*



*rotated left to make a legal 3-node*

# Insertion in a LLRB Tree



**Case 1.** Insert into a 2-node at the bottom.

- Do standard BST insert; color new link red
- If new red link is a right link, rotate left.

← to maintain symmetric order and perfect black balance




← to restore color invariants

insert C



corresponding red-black BST

right link red so rotate left



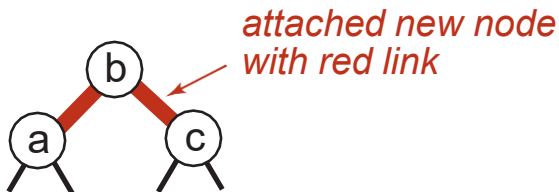
corresponding red-black BST

# Insertion in a LLRB Tree

Warmup 2. Insert into a tree with exactly 2 nodes.




larger

insert c





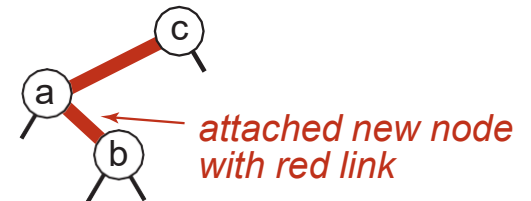
smaller

insert a



between

insert b



# Insertion in a LLRB Tree


**Case 2.** Insert into a 3-node at the bottom.

- Do standard BST insert; color new link red.
- Rotate to balance the 4-node (if needed).
- Flip colors to pass red link up one level.
- Rotate to make lean left (if needed).

← to maintain symmetric order and perfect black balance

← to restore color invariants

insert H



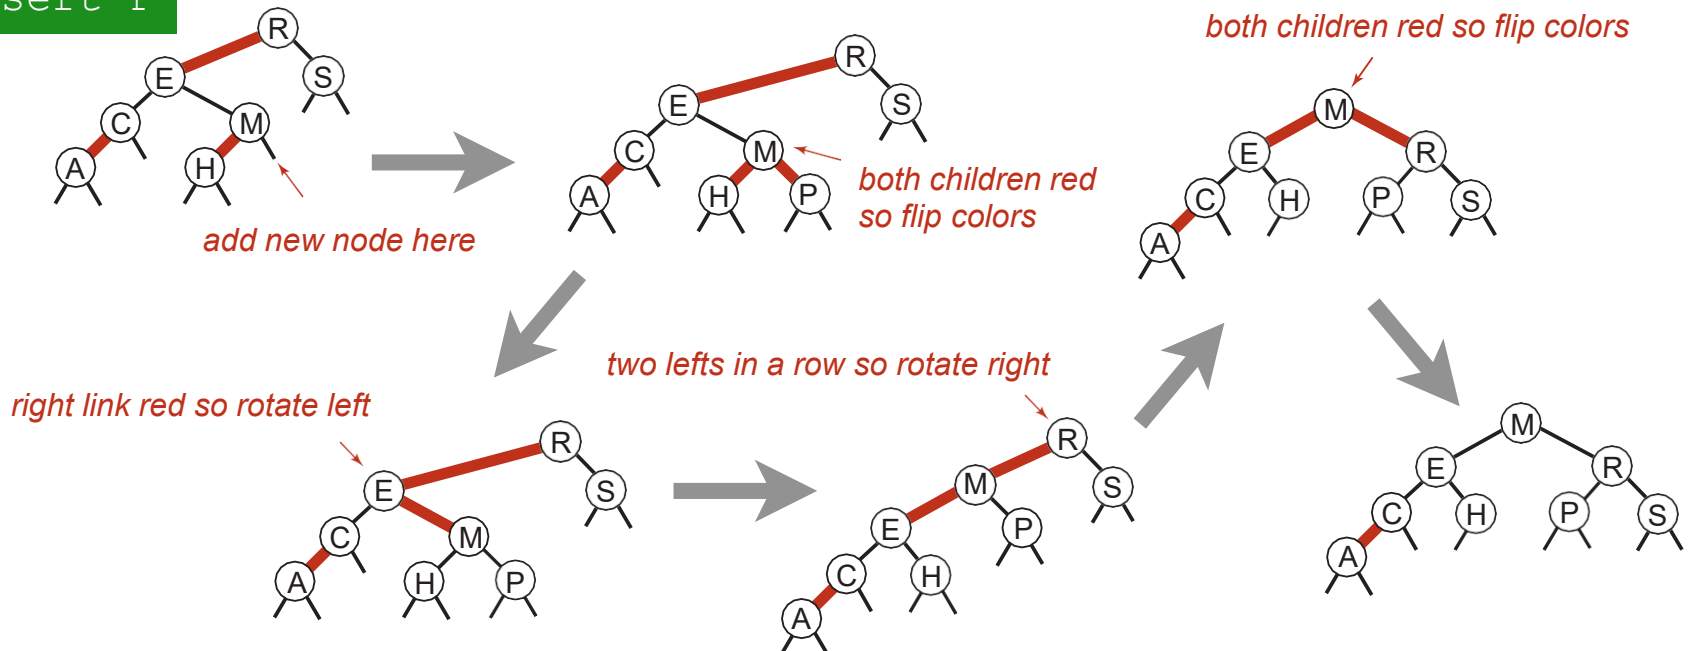
# Insertion in a LLRB Tree: Passing Red Links up the Tree

**Case 2.** Insert into a 3-node at the bottom.

- Do standard BST insert; color new link red.
- Rotate to balance the 4-node (if needed).
- Flip colors to pass red link up one level.
- Rotate to make lean left (if needed).
- **Repeat case 1 or case 2 up the tree (if needed).**

If the red link is passed up to a 3-node, we treat it the same way we do at the bottom

insert P



# LLRB Tree Construction Demo

insert S

insert E

insert A

insert R


insert C

insert H


insert X

insert M


two lefts in a row  
(rotate S right)



right link red  
(rotate A left)




both children red  
(flip colors)




two lefts in a row  
(rotate S right)


right link red  
(rotate E left)



both children red  
(flip colors)



right link red  
(rotate S left)










# LLRB Tree Construction Demo (Contd.)

```
insert P
```

```
insert L
```



# Insertion in a LLRB tree: Java implementation

Same code for all cases.

- Right child red, left child black: **rotate left.**
- Left child, left-left grandchild red: **rotate right.**
- Both children red: **flip colors.**


The standard BST code for insertion is in grey

```
private Node put(Node h, Key key, Value val)
{
    if (h == null) return new Node(key, val, RED);
    int cmp = key.compareTo(h.key);
    if (cmp < 0) h.left = put(h.left, key, val);
    else if (cmp > 0) h.right = put(h.right, key, val);
    else if (cmp == 0) h.val = val;

    if (isRed(h.right) && !isRed(h.left)) h = rotateLeft(h);
    if (isRed(h.left) && isRed(h.left.left)) h = rotateRight(h);
    if (isRed(h.left) && isRed(h.right)) flipColors(h);

    return h;
}
```

only a few extra lines of code provides near-perfect balance




insert at bottom  
(and color it red)

lean left


balance 4-node

split 4-node


# Insertion in a LLRB Tree: Visualization



255 insertions in descending order



255 insertions in ascending order



255 random insertions

- $N \rightarrow$  number of nodes
- max  $\rightarrow$  max number of nodes in a path
- avg  $\rightarrow$  average length of a path
- opt  $\rightarrow$  minimal length of a path

Height of tree is  $\leq 2 \lg N$  in the worst case.

- Every path from root to null link has same number of black links.
- Never two red links in-a-row.

Height of tree is  $\sim 1.00 \lg N$  in typical applications.

	Worst case			Average case		
	Search	Insert	Delete	Search	Insert	Delete
BST	$O(N)$	constants depend upon implementation		$O(1.39 \lg N)$		
2-3 tree	$O(c \lg N)$			$O(c \lg N)$		
LLRB	$O(2 \lg N)$			$O(1.00 \lg N)$		

exact value of coefficient unknown but extremely close to 1

# Additional Resources

- Left-leaning Red–Black Trees
  - <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.139.282>
- B-tree
  - <https://www.geeksforgeeks.org/b-tree-set-1-introduction-2/>

# Lecture Goals

- Develop balanced search trees with guaranteed logarithmic performance for search and insert (and many other operations).
- We begin with 2-3 trees, which are easy to analyze but hard to implement.
- Next, we consider red-black binary search trees, which we view as a novel way to implement 2-3 trees as binary search trees.
- Finally, we introduce **B-trees**, a generalization of 2-3 trees that are widely used to implement file systems.


# Motivation for B-tree: File system model

**Page.** Contiguous block of data (e.g., a file or 4,096-byte chunk).

**Probe.** First access to a page (e.g., from disk to memory).



slow



fast

**Property.** Time required for a probe is much larger than time to access data within a page.

**Cost model.** Number of probes.


**Goal.** Access data using minimum number of probes.

# B-trees (Bayer-McCreight, 1972)

Generalize 2-3 trees by allowing up to  $M - 1$  key-link pairs per node.

- At least 2 key-link pairs at root.
- At least  $M / 2$  key-link pairs in other nodes.
- External nodes contain client keys.
- Internal nodes contain copies of keys to guide search.

Choose  $M$  as large as possible so that  $M$  links fit in a page, e.g.,  $M = 1024$




Anatomy of a B-tree set (M = 6)

# Searching in a B-tree

- Start at root.
- Find interval for search key and take corresponding link.
- Search terminates in external node.

searching for E




Searching in a B-tree set ( $M = 6$ )



# Insertion in a B-tree

- Search for new key.
- Insert at bottom.
- Split nodes with M key-link pairs on the way up the tree.

inserting A




Inserting a new key into a B-tree set

# Balance in B-tree

**Proposition.** A search or an insertion in a B-tree of order  $M$  with  $N$  keys requires between  $\log_{M-1} N$  and  $\log_{M/2} N$  probes.

**Pf.** All internal nodes (besides root) have between  $M / 2$  and  $M - 1$  links.

**In practice.** Number of probes is at most 4.   $M = 1024; N = 62 \text{ billion}$   
 $\log_{M/2} N \leq 4$

**Optimization.** Always keep root page in memory.

# Balanced Trees in the Wild

- Red–black trees are widely used as system symbol tables.
  - Java: `java.util.TreeMap`, `java.util.TreeSet`.
  - C++ STL: `map`, `multimap`, `multiset`.
  - Linux kernel: completely fair scheduler, `linux/rbtree.h`.
  - Emacs: conservative stack scanning.
- B-tree variants. B+ tree, B\*tree, B# tree, ...
- B-trees (and variants) are widely used for file systems and databases.
  - Windows: NTFS.
  - Mac: HFS, HFS+.
  - Linux: ReiserFS, XFS, Ext3FS, JFS, BTRFS.
  - Databases: ORACLE, DB2, INGRES, SQL, PostgreSQL.