

Lecture 11

Shortest Paths

Jianchen Shan
Department of Computer Science
Hofstra University

Lecture Goals

- In this lecture we study **shortest-paths** problems. We begin by analyzing some basic properties of shortest paths and a generic algorithm for the problem.
- We introduce and analyze **Dijkstra's algorithm** for shortest-paths problems with nonnegative weights.
- Next, we consider an even faster **algorithm for DAGs**, which works even if the weights are negative.
- We conclude with the **Bellman–Ford–Moore** algorithm for edge-weighted digraphs with no negative cycles.
- We also consider **applications** ranging from content-aware fill to arbitrage.

Shortest Paths in an Edge-weighted Digraph

Given an edge-weighted digraph, find the shortest path from s to t .

edge-weighted digraph

4→5	0.35
5→4	0.35
4→7	0.37
5→7	0.28
7→5	0.28
5→1	0.32
0→4	0.38
0→2	0.26
7→3	0.39
1→3	0.29
2→7	0.34
6→2	0.40
3→6	0.52
6→0	0.58
6→4	0.93



shortest path from 0 to 6

0→2	0.26
2→7	0.34
7→3	0.39
3→6	0.52

Variants

❖ Which vertices?

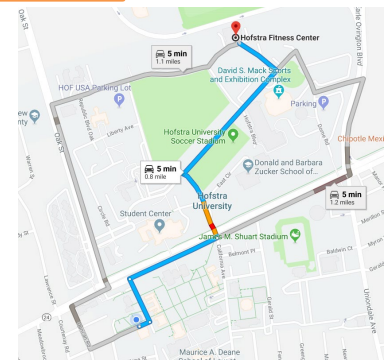
- Single source: from one vertex s to every other vertex.
- Source-sink: from one vertex s to another t .
- All pairs: between all pairs of vertices.

❖ Nonnegative weights?

❖ Cycles?

- Negative cycles.

Can we use BFS?



Simplifying assumption: Each vertex is reachable from s .

Weighted Directed Edge API

```
public class DirectedEdge
```

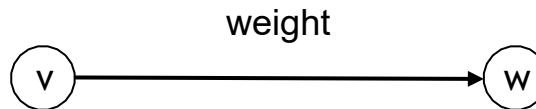
```
    DirectedEdge(int v, int w, double weight) //weighted edge v->w
```

```
    int from() // vertex v
```

```
    int to() // vertex w
```

```
    double weight() // the weight
```

```
    String toString() // string representation
```



Idiom for processing an edge *e*: `int v = e.from(), w = e.to();`

Weighted Edge: Java Implementation

```
public class DirectedEdge
{
    private final int v, w;
    private final double weight;

    public DirectedEdge(int v, int w, double weight)
    {
        this.v = v;
        this.w = w;
        this.weight = weight;
    }

    public int from()
    { return v; }

    public int to()
    { return w; }

    public int weight()
    { return weight; }
}
```

Edge-Weighted Graph API

```
public class EdgeWeightedDigraph
```

```
    EdgeWeightedDigraph(int V) // edge-weighted digraph with V vertices
```

```
    void addEdge(DirectedEdge e) // add weighted directed edge e
```

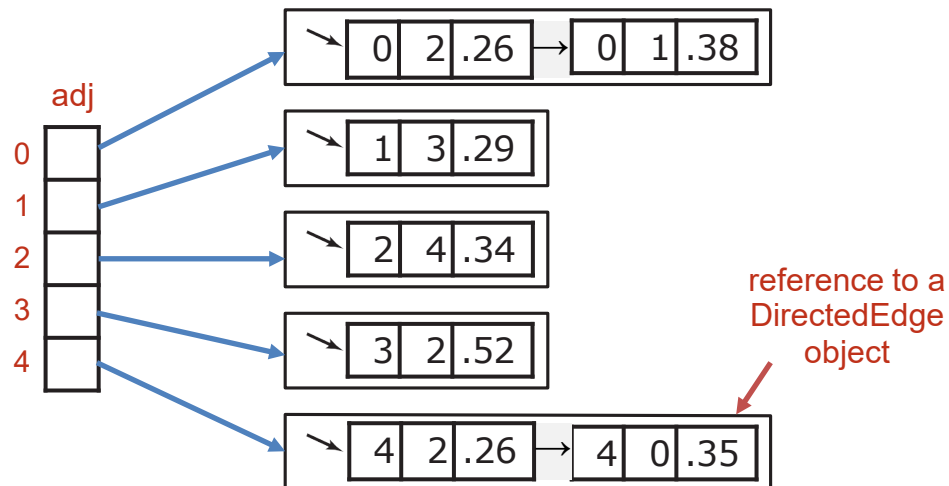
```
    Iterable<DirectedEdge> adj(int v) // edges pointing from v
```

```
    Iterable<DirectedEdge> edges() // all edges in this graph
```

```
    int V() // number of vertices
```

```
    int E() // number of edges
```

```
    String toString() // string representation
```



Edge-Weighted Digraph: Adjacency-Lists Implementation

```
public class EdgeWeightedDigraph
{
    private final int V;
    private final List<DirectedEdge>[] adj;

    public EdgeWeightedDigraph (int V)
    {
        this.V = V;
        adj = (List<DirectedEdge>[]) new ArrayList[V];
        for (int v = 0; v < V; v++)
            adj[v] = new ArrayList<DirectedEdge>();
    }

    public void addEdge(DirectedEdge e)
    {
        int v = e.from();
        adj[v].add(e);
    }

    public Iterable<DirectedEdge> adj(int v)
    {
        return adj[v];
    }
}
```

← add edge $e = v \rightarrow w$ to
only v 's adjacency lists

Single-source Shortest Paths API

Goal. Find the shortest path from s to every other vertex.

```
public class SP
```

```
    SP(EdgeWeightedGraph G, int s) // shortest paths from s in graph G
```

```
    double distTo(int v) // length of shortest path from s to v
```

```
    Iterable<DirectedEdge> pathTo(int v) // shortest path from s to v
```

```
SP sp = new SP(G, s);
for (int v = 0; v < G.V(); v++)
{
    StdOut.printf("%d to %d (0.2f): ", s, v, sp.distTo(v));
    for (DirectedEdge e : sp.pathTo(v)) {
        StdOut.print(e + " ");
        StdOut.println();
    }
}
```

```
0 to 0 (0.00):
0 to 1 (1.05): 0->4 0.38 4->5 0.35 5->1 0.32
0 to 2 (0.26): 0->2 0.26
0 to 3 (0.99): 0->2 0.26 2->7 0.34 7->3 0.39
0 to 4 (0.38): 0->4 0.38
0 to 5 (0.73): 0->4 0.38 4->5 0.35
0 to 6 (1.51): 0->2 0.26 2->7 0.34 7->3 0.39 3->6 0.52
0 to 7 (0.60): 0->2 0.26 2->7 0.34
```

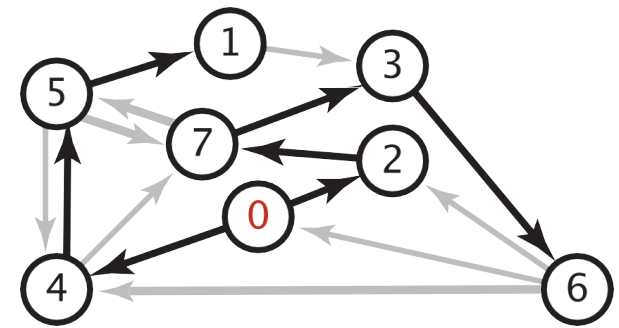

Data Structures for Single-source Shortest Paths

Goal. Find the shortest path from s to every other vertex.

Observation. A shortest-paths tree (SPT) solution exists.

Consequence. Can represent the SPT with two vertex-indexed arrays:

- $\text{distTo}[v]$ is length of shortest path from s to v .
- $\text{edgeTo}[v]$ is last edge on shortest path from s to v .



shortest-paths tree from 0

	edgeTo[]	distTo[]
0	null	0
1	5->1 0.32	1.05
2	0->2 0.26	0.26
3	7->3 0.37	0.97
4	0->4 0.38	0.38
5	4->5 0.35	0.73
6	3->6 0.52	1.49
7	2->7 0.34	0.60

parent-link representation

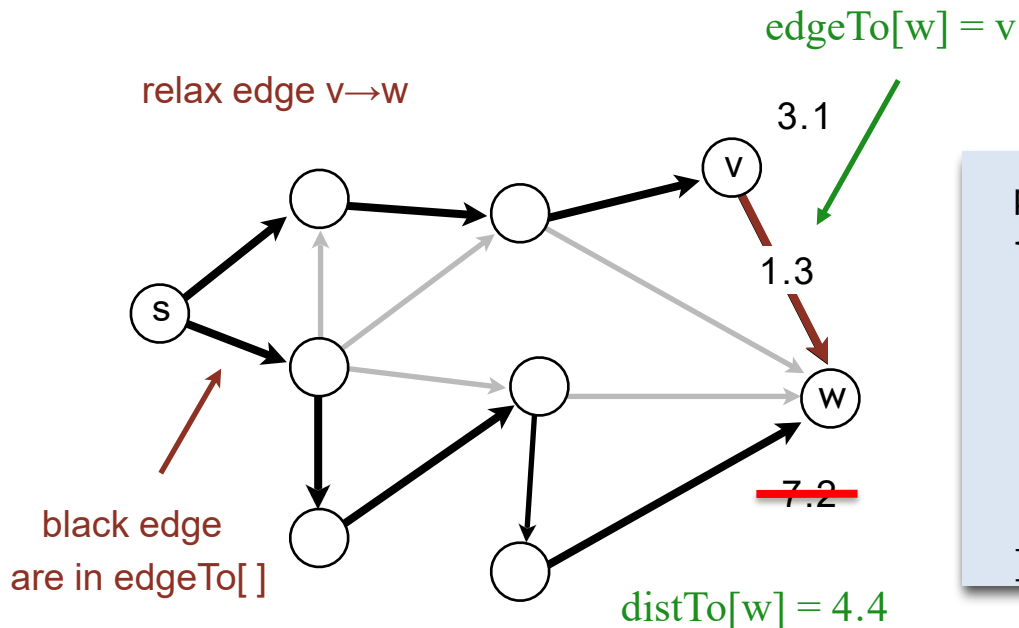
```
public double distTo(int v)
{ return distTo[v]; }

public Iterable<DirectedEdge> pathTo(int v)
{
    Stack<DirectedEdge> path = new Stack<DirectedEdge>();
    for (DirectedEdge e = edgeTo[v]; e != null; e = edgeTo[e.from()])
        path.push(e);
    return path;
}
```

Edge Relaxation

Relax edge $e = v \rightarrow w$. (basic of building SPT)

- $\text{distTo}[v]$ is length of shortest **known** path from s to v .
- $\text{distTo}[w]$ is length of shortest **known** path from s to w .
- $\text{edgeTo}[w]$ is last edge on shortest **known** path from s to w .
- If $e = v \rightarrow w$ gives shorter path to w through v , update $\text{distTo}[w]$ and $\text{edgeTo}[w]$.



```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] + e.weight();
        edgeTo[w] = e;
    }
}
```

Shortest-paths Optimality conditions

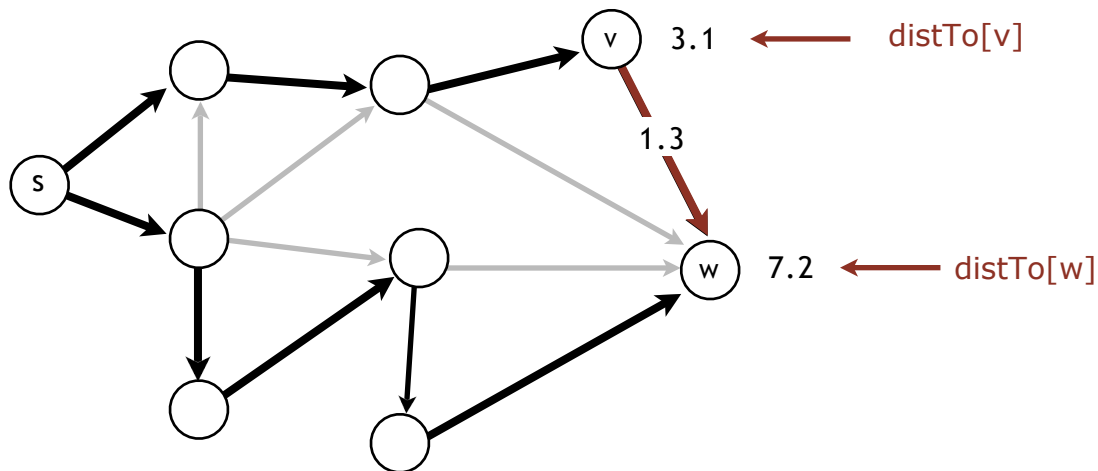
Proposition. Let G be an edge-weighted digraph.

Then $\text{distTo}[]$ are the shortest path distances from s *iff*:

- $\text{distTo}[s] = 0$.
- For each vertex v , $\text{distTo}[v]$ is the length of some path from s to v .
- For each edge $e = v \rightarrow w$, $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.

Pf. \Rightarrow [necessary]

- Suppose that $\text{distTo}[w] > \text{distTo}[v] + e.\text{weight}()$ for some edge $e = v \rightarrow w$.
- Then, e gives a path from s to w (through v) of length less than $\text{distTo}[w]$.



Shortest-paths Optimality conditions

Proposition. Let G be an edge-weighted digraph.

Then $\text{distTo}[]$ are the shortest path distances from s *iff*:

- $\text{distTo}[s] = 0$.
- For each vertex v , $\text{distTo}[v]$ is the length of some path from s to v .
- For each edge $e = v \rightarrow w$, $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.

Pf. \Leftarrow [sufficient]

- Suppose that $s = v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k = w$ is a shortest path from s to w .

- Then,

$$\text{distTo}[v_1] \leq \text{distTo}[v_0] + e_1.\text{weight}()$$

$$\text{distTo}[v_2] \leq \text{distTo}[v_1] + e_2.\text{weight}()$$

.....

$$\text{distTo}[v_k] \leq \text{distTo}[v_{k-1}] + e_k.\text{weight}()$$



$e_i = i^{\text{th}}$ edge on shortest path from s to w

- Add inequalities; simplify; and substitute $\text{distTo}[v_0] = \text{distTo}[s] = 0$:

$$\text{distTo}[w] = \text{distTo}[v_k] \leq \underline{e_1.\text{weight}() + e_2.\text{weight}() + \dots + e_k.\text{weight}()}$$

- Thus, $\text{distTo}[w]$ is the weight of shortest path to w .



weight of shortest path from s to w

Generic Shortest-paths Algorithm

Generic algorithm (to compute SPT from s)

For each vertex v : $\text{distTo}[v] = \infty$.

For each vertex v : $\text{edgeTo}[v] = \text{null}$.

$\text{distTo}[s] = 0$.

Repeat until done:

- Relax any edge.

Proposition. Generic algorithm computes SPT (if it exists) from s .

Pf.

- Throughout algorithm, $\text{distTo}[v]$ is the length of a simple path from s to v (and $\text{edgeTo}[v]$ is last edge on path).
- Each successful relaxation decreases $\text{distTo}[v]$ for some v .
- The entry $\text{distTo}[v]$ can decrease at most a finite number of times.


Efficient implementations. How to choose which edge to relax?

- Ex 1. Dijkstra's algorithm. (**nonnegative weights, directed cycles**).
- Ex 2. Topological sort algorithm. (**no directed cycles**).
- Ex 3. Bellman–Ford algorithm. (**no neigitive cycles**).

Dijkstra's Algorithm

- Consider vertices in increasing order of distance from s
 - (non-tree vertex with the lowest distTo[] value).
- Add vertex to tree and relax all edges pointing from that vertex.

choose vertex 5
 relax all edges adjacent from 5
 choose vertex 2
 relax all edges adjacent from 2
 choose vertex 3
 relax all edges adjacent from 3
 choose vertex 6
 relax all edges adjacent from 6



choose source vertex 0
 relax all edges adjacent from 0
 choose vertex 1
 relax all edges adjacent from 1

choose vertex 7
 relax all edges adjacent from 7
 choose vertex 4
 relax all edges adjacent from 4



v distTo[]			
→ 0	∞	0	
→ 1	∞	5	
→ 2	∞	17	15 14
→ 3	∞	20	17
→ 4	∞	9	
→ 5	∞	14	13
→ 6	∞	29	26 25
→ 7	∞	8	

v edgeTo[]			
0	-		
1	-	0	
2	-	1	7 5
3	-	1	2
4	-	0	
5	-	7	4
6	-	4	5 2
7	-	0	

Dijkstra's Algorithm: Correctness Proof

Proposition. Dijkstra's algorithm computes a SPT in any edge-weighted digraph with nonnegative weights.

Pf.

- Each edge $e = v \rightarrow w$ is relaxed exactly once (when v is relaxed),
 - leaving $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.
- Inequality holds until algorithm terminates because:
 - $\text{distTo}[w]$ cannot increase  $\text{distTo}[\]$ values are monotone decreasing
 - $\text{distTo}[v]$ will not change  we choose lowest $\text{distTo}[\]$ value at each step (and edge weights are nonnegative)
- Thus, upon termination, shortest-paths optimality conditions hold.

Dijkstra's Algorithm: Java Implementation

```
public class DijkstraSP
{
    private DirectedEdge [] edgeTo;
    private double [] distTo;
    private MinPQ<Double> pq;

    public DijkstraSP(EdgeWeightedDigraph G, int s)
    {
        edgeTo = new DirectedEdge[G.V()];
        distTo = new double[G.V()];
        pq = new MinPQ<Double>(G.V());

        for (int v = 0; v < G.V(); v++)
            distTo[v] = Double.POSITIVE_INFINITY;
        distTo[s] = 0.0;

        pq.insert(s, 0.0);
        while (!pq.isEmpty())
        {
            int v = pq.delMin();
            for (DirectedEdge e : G.adj(v))
                relax(e);
        }
    }
}
```

← relax vertices in order
of distance from s

Dijkstra's Algorithm: Java Implementation

```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] + e.weight();
        edgeTo[w] = e;
        if (pq.contains(w)) pq.decreaseKey (w, distTo[w]);
        else                pq.insert      (w, distTo[w]);
    }
}
```

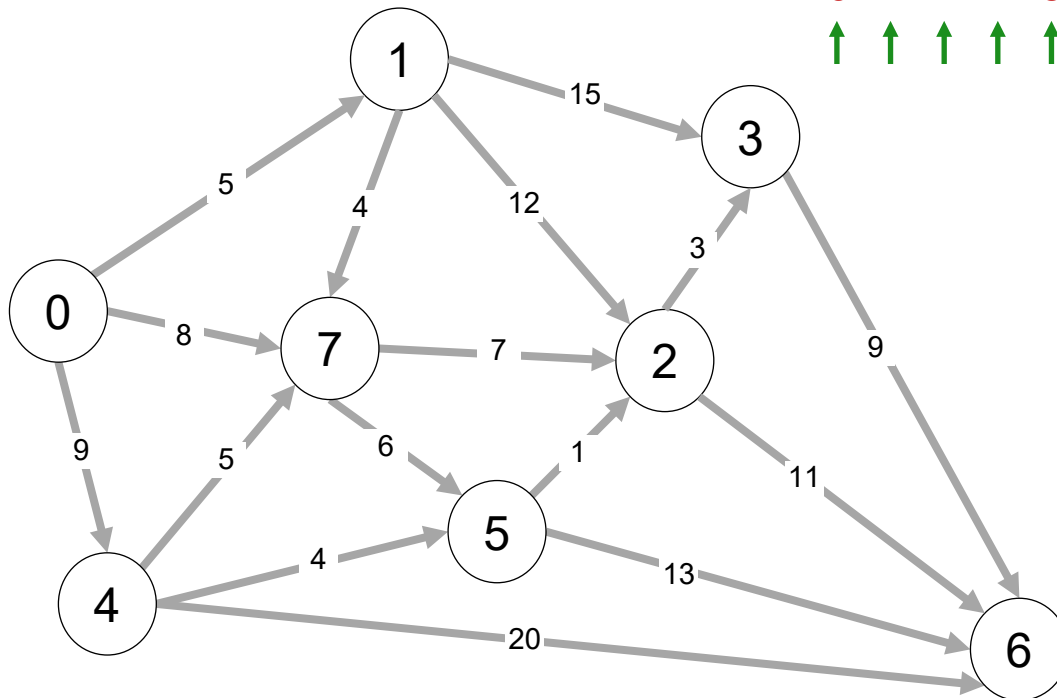
← update PQ

Shortest Paths in Edge-weighted DAG

Suppose that an edge-weighted digraph has no directed cycles. Is it easier to find shortest paths than in a general digraph?

Yes!

- Consider vertices in topological order.
- Relax all edges pointing from that vertex



0 1 4 7 5 2 3 6
↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑

v	distTo[]	
0	∞	0
1	∞	5
2	∞	17 15 14
3	∞	20 17
4	∞	9
5	∞	13
6	∞	29 26 25
7	∞	8

v	edgeTo[]	
0	-	
1	-	0
2	-	1 7 5
3	-	1 2
4	-	0
5	-	4
6	-	4 5 2
7	-	0

Shortest Paths in Edge-weighted DAG: Correctness Proof

Proposition. Topological sort algorithm computes SPT in any edge-weighted DAG in time proportional to $E + V$.

edge weights
can be negative!

Pf.

- Each edge $e = v \rightarrow w$ is relaxed exactly once (when v is relaxed),
 - leaving $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.
- Inequality holds until algorithm terminates because:
 - $\text{distTo}[w]$ cannot increase ← $\text{distTo}[\]$ values are monotone decreasing
 - $\text{distTo}[v]$ will not change ← because of topological order, no edge pointing to v will be relaxed after v is relaxed
- Thus, upon termination, shortest-paths optimality conditions hold.

Shortest Paths in Edge-weighted DAG: Java Implementation

```
public class AcyclicSP
{
    private DirectedEdge[] edgeTo;
    private double[] distTo;

    public AcyclicSP (EdgeWeightedDigraph G, int s)
    {
        edgeTo = new DirectedEdge[G.V()];
        distTo = new double[G.V()];

        for (int v = 0; v < G.V(); v++)
            distTo[v] = Double.POSITIVE_INFINITY;
        distTo[s] = 0.0;

        Topological topological = new Topological(G);
        for (int v : topological.order())
            for (DirectedEdge e : G.adj(v))
                relax(e);
    }
}
```

← topological order

Longest Paths in Edge-weighted DAG


Formulate as a shortest paths problem in edge-weighted DAGs.

- Negate all weights.
 - Find shortest paths.
 - Negate weights in result.
- equivalent: reverse sense of equality in relax()

longest paths input

shortest paths input


5->4	0.35	5->4	-0.35
4->7	0.37	4->7	-0.37
5->7	0.28	5->7	-0.28
5->1	0.32	5->1	-0.32
4->0	0.38	4->0	-0.38
0->2	0.26	0->2	-0.26
3->7	0.39	3->7	-0.39
1->3	0.29	1->3	-0.29
7->2	0.34	7->2	-0.34
6->2	0.40	6->2	-0.40
3->6	0.52	3->6	-0.52
6->0	0.58	6->0	-0.58
6->4	0.93	6->4	-0.93



Key point. Topological sort algorithm works even with negative weights.

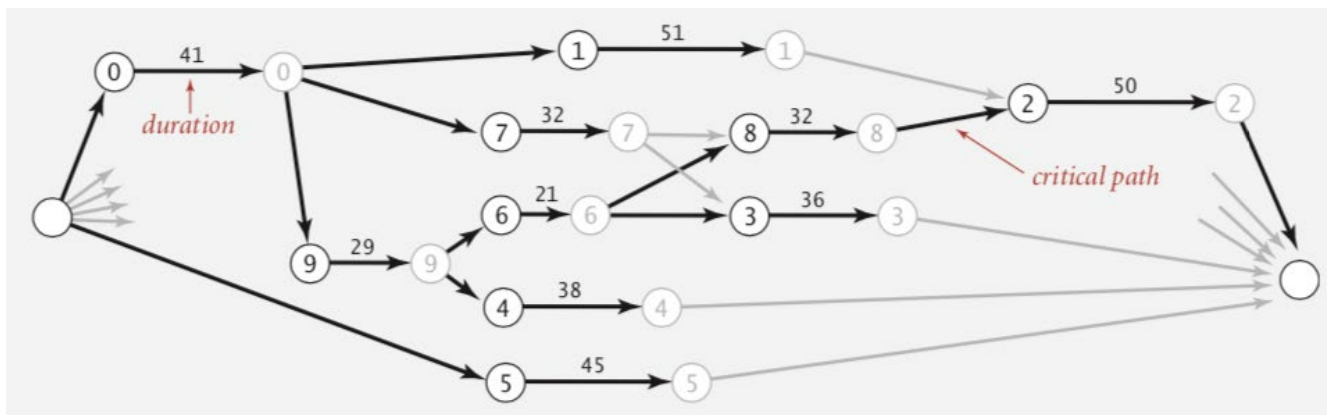
Longest Paths in Edge-weighted DAG: Application

Parallel job scheduling. Given a set of jobs with durations and precedence constraints, schedule the jobs (by finding a start time for each) so as to achieve the minimum completion time, while respecting the constraints.





job	duration	must complete before		
0	41.0	1	7	9
1	51.0	2		
2	50.0			
3	36.0			
4	38.0			
5	45.0			
6	21.0	3	8	
7	32.0	3	8	
8	32.0	2		
9	29.0	4	6	

Use **longest path** from the source to schedule each job.



Shortest Paths with Negative weights

Dijkstra. Doesn't work with negative edge weights.




Conclusion.
Need a different algorithm.

Dijkstra selects vertex 3 immediately after 0.
But shortest path from 0 to 3 is $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$.

Adding 9 to each edge weight changes the
shortest path from $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$ to $0 \rightarrow 3$.

Re-weighting. Add a constant to every edge weight doesn't work.

- A **negative cycle** is a directed cycle whose sum of edge weights is negative.
- A SPT exists **iff** no negative cycles, assuming all vertices reachable from s



negative cycle $(-0.66 + 0.37 + 0.28)$

$5 \rightarrow 4 \rightarrow 7 \rightarrow 5$

shortest path from 0 to 6

$0 \rightarrow 4 \rightarrow 7 \rightarrow 5 \rightarrow 4 \rightarrow 7 \rightarrow 5 \dots \rightarrow 1 \rightarrow 3 \rightarrow 6$

4->5	0.35
5->4	-0.66
4->7	0.37
5->7	0.28
7->5	0.28
5->1	0.32
0->4	0.38
0->2	0.26
7->3	0.39
1->3	0.29
2->7	0.34
6->2	0.40
3->6	0.52
6->0	0.58
6->4	0.93

Bellman-Ford Algorithm

Bellman-Ford algorithm

For each vertex v : $\text{distTo}[v] = \infty$.

For each vertex v : $\text{edgeTo}[v] = \text{null}$.

$\text{distTo}[s] = 0$.

Repeat $V-1$ times:


- Relax each edge.

```
for (int i = 1; i < G.V(); i++)  
    for (int v = 0; v < G.V(); v++)  
        for (DirectedEdge e : G.adj(v))  
            relax(e);
```

← pass i (relax each edge)

Bellman-Ford Algorithm

Repeat $V - 1$ times: relax all E edges.



v	distTo[]		
0	∞	0	
1	∞	5	
2	∞	17	14
3	∞	20	17
4	∞	9	
5	∞	13	
6	∞	28	26 25
7	∞	8	

v	edgeTo[]		
0	-		
1	-	0	
2	-	1	5
3	-	1	2
4	-	0	
5	-	4	
6	-	2	5 2
7	-	0	

pass 1 pass 2 pass 3 (no further changes) pass 4-7 (no further changes)

0→1 0→4 0→7 1→2 1→3 1→7 2→3 2→6 3→6 4→5 4→6 4→7 5→2 5→6 7→2 7→5

Bellman–Ford Algorithm: Correctness Proof

- **Proposition.** Let $s = v_0 \rightarrow v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k = v$ be a shortest path from s to v . Then, after pass i , $\text{distTo}[v_i] = d^*(v_i)$.

length of shortest path from s to v_i

- **Pf.** [by induction on i]

- Inductive hypothesis: after pass i , $\text{distTo}[v_i] = d^*(v_i)$.
- Since $\text{distTo}[v_{i+1}]$ is the length of some path from s to v_{i+1} , we must have $\text{distTo}[v_{i+1}] \geq d^*(v_{i+1})$.
- Immediately after relaxing edge $v_i \rightarrow v_{i+1}$ in pass $i+1$, we have




$$\begin{aligned}\text{distTo}[v_{i+1}] &\leq \text{distTo}[v_i] + \text{weight}(v_i, v_{i+1}) \\ &= d^*(v_i) + \text{weight}(v_i, v_{i+1}) \\ &= d^*(v_{i+1}).\end{aligned}$$


- Thus, at the end of pass $i+1$, $\text{distTo}[v_{i+1}] = d^*(v_{i+1})$.
- **Corollary.** Bellman–Ford computes shortest path distances.
- **Pf.** There exists a shortest path from s to v with at most $V - 1$ edges.
 $\Rightarrow \leq V - 1$ passes.

Bellman-Ford Algorithm Visualization

passes 4



7



10



13



SPT




Bellman-Ford Algorithm Analysis

Observation. If $\text{distTo}[v]$ does not change during pass i , no need to relax any edge pointing from v in pass $i + 1$.

Queue-based implementation of Bellman-Ford. Maintain **queue** of vertices whose $\text{distTo}[\]$ values needs updating.

In the worst case, the running time is still proportional to $E \times V$. But much faster in practice.



be careful to keep at most one copy of each vertex on queue

Observation. If there is a negative cycle, Bellman-Ford gets stuck in loop, updating $\text{distTo}[\]$ and $\text{edgeTo}[\]$ entries of vertices in the cycle.

Proposition. If any vertex v is updated in phase V , there exists a negative cycle (and can trace back $\text{edgeTo}[v]$ entries to find it).


Finding a negative cycle

Negative Cycle Application: Arbitrage Detection

Problem. Given table of exchange rates, is there an arbitrage opportunity?

	USD	EUR	GBP	CHF	CAD
USD	1	0.741	0.657	1.061	1.011
EUR	1.350	1	0.888	1.433	1.366
GBP	1.521	1.126	1	1.614	1.538
CHF	0.943	0.698	0.620	1	0.953
CAD	0.995	0.732	0.650	1.049	1

Ex. \$1,000 \Rightarrow 741 Euros \Rightarrow 1,012.206 Canadian dollars \Rightarrow \$1,007.14497.

$$1000 \times 0.741 \times 1.366 \times 0.995 = 1007.14497$$



Negative Cycle Application: Arbitrage Detection

Currency exchange graph.

- Vertex = currency.
- Edge = transaction, with weight equal to exchange rate.
- Find a directed cycle whose product of edge weights is > 1 .


Challenge. Express as a negative cycle detection problem.

$$0.741 * 1.366 * .995 = 1.00714497$$



$$\begin{aligned} & -\ln(.741) \quad -\ln(1.366) \quad -\ln(.995) \\ & \quad \downarrow \quad \downarrow \quad \downarrow \\ & .2998 \quad -.3119 \quad + .0050 \\ & = -.0071 \end{aligned}$$

Remark. Fastest algorithm is extraordinarily valuable!



Model as a negative cycle detection problem by taking logs.

- Let weight of edge $v \rightarrow w$ be $-\ln(\text{exchange rate from currency } v \text{ to } w)$.
- Multiplication turns to addition; > 1 turns to < 0 .
- Find a directed cycle whose sum of edge weights is < 0 (negative cycle).

Single Source Shortest-paths Implementation: Cost Summary


algorithm	restriction	typical case	worst case	extra space
topological sort	no directed cycles	$E + V$	$E + V$	V
Dijkstra (binary heap)	no negative weights	$E \log V$	$E \log V$	V
Bellman-Ford	no negative cycles	$E V$	$E V$	V
Bellman-Ford (queue-based)		$E + V$	$E V$	V

- **Remark 1.** Directed cycles make the problem harder.
- **Remark 2.** Negative weights make the problem harder.
- **Remark 3.** Negative cycles makes the problem intractable.

Backup Slides

Dijkstra's Algorithm Analysis

- Dijkstra's algorithm seem familiar?
 - Prim's algorithm is essentially the same algorithm.
 - Both are in a family of algorithms that compute a graph's spanning tree.
- Main distinction: Rule used to choose next vertex for the tree.
 - Prim's: Closest vertex to the tree (via an undirected edge).
 - Dijkstra's: Closest vertex to the source (via a directed path).
- Note: DFS and BFS are also in this family of algorithms.



$O(E \log V)$

operation	frequency	time per op
Insert	E	$\log V$
delete min	E	$\log V$
decrease key	E	$\log V$

Computing spanning trees in graphs