

Lecture 11

Shortest Paths

Department of Computer Science
Hofstra University

Lecture Goals


- In this lecture we study shortest-paths problems. We begin by analyzing some basic properties of shortest paths and a generic algorithm for the problem.
- For single-source shortest path, we consider:
 - Dijkstra's algorithm
 - Bellman–Ford algorithm
 - Topological Sort for DAG
- For all-pairs shortest path, we conclude:
 - Floyd Warshall Algorithm
 - Johnson's Algorithm

Shortest Paths in an Edge-weighted Digraph

Given an edge-weighted digraph, find the shortest path from source vertex s to t .

edge-weighted digraph

4→5	0.35
5→4	0.35
4→7	0.37
5→7	0.28
7→5	0.28
5→1	0.32
0→4	0.38
0→2	0.26
7→3	0.39
1→3	0.29
2→7	0.34
6→2	0.40
3→6	0.52
6→0	0.58
6→4	0.93



shortest path from 0 to 6

0→2	0.26
2→7	0.34
7→3	0.39
3→6	0.52

Variants


❖ Which vertices?

- Single source: from source vertex s to every other vertex.
- Source-sink: from source vertex s to another t .
- All pairs: between all pairs of vertices.

❖ Nonnegative weights?

❖ Cycles?

- Negative cycles.



Simplifying assumption: Each vertex is reachable from s .

Shortest Paths in an Unweighted Digraph


- BFS (Breadth-First Search) can find shortest paths in unweighted graphs.
 - BFS visits nodes in order of their distance from the source vertex, ensuring the first path found to any node is the shortest possible path in terms of the number of edges.
 - Time complexity: $O(V+E)$
- Advantages:
 - Optimal for unweighted graphs
 - Simpler implementation than Dijkstra's
- Limitations:
 - Only works for unweighted graphs
 - Not suitable for graphs with negative edges

Edge Relaxation

Relax edge $e = u \rightarrow v$ with weight $w(u,v)$. (We also write uv to denote $u \rightarrow v$)

- $\text{distTo}[u]$ is length of shortest **known** path from s to u .
- $\text{distTo}[v]$ is length of shortest **known** path from s to v .
- $\text{prevNode}[v]$ is the previous vertex on shortest **known** path from s to v .
- If $e = u \rightarrow v$ gives shorter path to v through u , update $\text{distTo}[v]$ and $\text{prevNode}[v]$.
 - $\text{distTo}[v] = \min(\text{distTo}[v], \text{distTo}[u] + w(u,v)); \text{prevNode}[v] = u$

After relaxing edge uv , the shortest path from s to v is updated to go through vertex u , with cost of 4.4



Previous shortest path from s to v goes through vertex x , with cost of 7.2

```
private void relax(DirectedEdge e)
{
    Int u = e.from(), v = e.to();
    if (distTo[v] > distTo[u] + w(u,v))
    {
        distTo[v] = distTo[u] + w(u,v);
        prevNode[v] = u;
    }
}
```

OLD $\text{distTo}[v] = 7.2 > \text{distTo}[u] + w(u,v) = 3.1 + 1.3 = 4.4$
NEW $\text{distTo}[v] \leftarrow \text{distTo}[u] + w(u,v) = 4.4$,
 $\text{prevNode}[v] = u$

Generic Shortest-paths Algorithm

Generic algorithm (to compute SPT from s)

For each vertex v : $\text{distTo}[v] = \infty$.

For each vertex v : $\text{prevNode}[v] = \text{null}$.

$\text{distTo}[s] = 0$.

Repeat until done:

- Relax any edge.

Proposition. Generic algorithm computes SPT (if it exists) from s .

Pf.

- Throughout algorithm, $\text{distTo}[v]$ is the length of a simple path from s to v (and $\text{prevNode}[v]$ is its previous vertex on the path).
- Each successful relaxation decreases $\text{distTo}[v]$ for some v .
- The entry $\text{distTo}[v]$ can decrease at most a finite number of times.

Efficient implementations. How to choose which edge to relax?

- Ex 1. Dijkstra's algorithm. (**no negative weights**).
- Ex 2. Bellman–Ford algorithm. (**negative weights, can detect negative cycles**).
- Ex 3. Topological sort. (**DAG with no directed cycles**)



Dijkstra's Algorithm

- Initialization:
 - Set the distance to the source vertex as 0 and to all other vertices as infinity.
 - Mark all vertices as unvisited and store them in a priority queue.
- Main Loop:
 - Visit the **unvisited vertex u** with **the shortest known distance** from the queue.
 - For each **unvisited neighbor vertex v** of vertex u, calculate its tentative distance through the current vertex. **If this distance is smaller than the previously recorded distance, update it with edge relaxation for edge uv.**
 - Mark the current vertex as visited once all its neighbors are processed.
- Termination:
 - The algorithm continues until all reachable vertices are visited.
- Time complexity: $O(V \log V + V)$ for Binary Heap implementation
- Notes:
 - Dijkstra's Algorithm is greedy and optimal: any vertex that has been visited should have its shortest distance to the source.
 - It works for both undirected and directed graphs. The only difference is the function for getting the neighbors of vertex v, as each undirected edge is treated as two directed edges in opposite directions.)


Dijkstra's Algorithm: Correctness Proof

Proposition. Dijkstra's algorithm computes a SPT in any edge-weighted digraph with nonnegative weights.

Proof.

- Each edge $e = u \rightarrow v$ is relaxed exactly once (when vertex u is visited), afterwards:
 - $\text{distTo}[v] \leq \text{distTo}[u] + w(u,v)$.
- Inequality holds until algorithm terminates because:
 - $\text{distTo}[v]$ cannot increase  $\text{distTo}[\]$ values are monotone decreasing
 - $\text{distTo}[u]$ will not change  we choose lowest $\text{distTo}[\]$ value at each step (and edge weights are nonnegative)
- Thus, upon termination, shortest-paths optimality conditions hold.

Toy Example: find shortest path starting from source vertex S for undirected graph
SD: Shortest Distance. PN: Previous Node




N1	SD	PN
S	0	
A	∞	
B	∞	

Visit S




N1	SD	PN
S	0	
A	2	S
B	4	S

Visit A




N1	SD	PN
S	0	
A	2	S
B	3	A

Visit B




N1	SD	PN
S	0	
A	2	S
B	3	A




N1	SD	PN
S	0	
A	∞	
B	∞	

Visit S



N1	SD	PN
S	0	
A	2	S
B	4	S

Visit A




N1	SD	PN
S	0	
A	2	S
B	4	S

Visit B




N1	SD	PN
S	0	
A	2	S
B	4	S

Toy Example: find shortest path starting from source vertex S for directed graph




N1	SD	PN
S	0	
A	∞	
B	∞	

Visit S




N1	SD	PN
S	0	
A	2	S
B	4	S

Visit A




N1	SD	PN
S	0	
A	2	S
B	4	S

Visit B




N1	SD	PN
S	0	
A	2	S
B	4	S




N1	SD	PN
S	0	
A	∞	
B	∞	

Visit S




N1	SD	PN
S	0	
A	∞	
B	4	S

Visit B




N1	SD	PN
S	0	
A	5	B
B	4	S

Visit A




N1	SD	PN
S	0	
A	5	B
B	4	S

Example Graph



Initialize

2. Assign to all nodes a tentative distance value



Visited Nodes: []


Unvisited Nodes: [A, B, C, D, E, F]

Node	Shortest Distance	Previous Node
A	0	
B	∞	
C	∞	
D	∞	
E	∞	
F	∞	

Visit vertex A

3. For the current node calculate the distance to all unvisited neighbours

3.1. Update shortest distance, if new distance is shorter than old distance



Visited Nodes: []

Unvisited Nodes: [A, B, C, D, E, F]

Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	∞	
D	8	A
E	∞	
F	∞	

OLD $\text{distTo}[B] = \infty > \text{distTo}[A] + w(A,B) = 0+2 = 2$


NEW $\text{distTo}[B] \leftarrow \text{distTo}[A] + w(A,B) = 2$, $\text{prevNode}[B] = A$

OLD $\text{distTo}[D] = \infty > \text{distTo}[A] + w(A,D) = 0+8 = 8$

NEW $\text{distTo}[D] \leftarrow \text{distTo}[A] + w(A,D) = 8$, $\text{prevNode}[D] = A$

Visit vertex B

3. For the current node calculate the distance to all unvisited neighbours
3.1. Update shortest distance, if new distance is shorter than old distance



Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	∞	
D	7	B
E	8	B
F	∞	

OLD $\text{distTo}[D] = 8 > \text{distTo}[B] + w(B,D) = 2+5 = 7$


NEW $\text{distTo}[D] \leftarrow \text{distTo}[B] + w(B,D) = 7$, $\text{prevNode}[D] = B$

OLD $\text{distTo}[E] = \infty > \text{distTo}[B] + w(B,E) = 2+6 = 8$

NEW $\text{distTo}[E] \leftarrow \text{distTo}[B] + w(B,E) = 8$, $\text{prevNode}[E] = B$

Visit vertex D

3. For the current node calculate the distance to all unvisited neighbours
3.1. Update shortest distance, if new distance is shorter than old distance



Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	∞	
D	7	B
E	8	B
F	9	D

OLD $\text{distTo}[E] = 8 < \text{distTo}[D] + w(D,E) = 7+3 = 10$

No update, $\text{distTo}[E]$ stays 8, $\text{prevNode}[E]$ stays B


OLD $\text{distTo}[F] = \infty > \text{distTo}[D] + w(D,F) = 7+2 = 9$

NEW $\text{distTo}[F] \leftarrow \text{distTo}[D] + w(D,F) = 9$, $\text{prevNode}[F] = D$

Visit vertex E

3. For the current node calculate the distance to all unvisited neighbours

3.1. Update shortest distance, if new distance is shorter than old distance



Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	17	E
D	7	B
E	8	B
F	9	D

Visited Nodes: [A, B, D] Unvisited Nodes: [C, E, F]

OLD $\text{distTo}[C] = \infty > \text{distTo}[E] + w(E.C) = 8 + 9 = 17$

NEW $\text{distTo}[C] \leftarrow \text{distTo}[E] + w(E.C) = 17$, $\text{prevNode}[C] = E$


OLD $\text{distTo}[F] = 9 = \text{distTo}[E] + w(E.F) = 8 + 1 = 9$

No update, $\text{distTo}[F]$ stays 9, $\text{prevNode}[F] = D$ (You can also update $\text{prevNode}[F] = E$.)

Visit vertex F

3. For the current node calculate the distance to all unvisited neighbours

3.1. Update shortest distance, if new distance is shorter than old distance




Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	12	F
D	7	B
E	8	B
F	9	D

OLD $\text{distTo}[C] = 17 > \text{distTo}[F] + w(F,C) = 9 + 3 = 12$

NEW $\text{distTo}[C] \leftarrow \text{distTo}[F] + w(F,C) = 12$, $\text{prevNode}[C] = F$

Visit vertex C




Nothing changes, since C has no unvisited neighbor vertices

End of Algorithm

- Table contains the shortest distance to each vertex N from the source vertex A, and its previous vertex in the shortest path

4. Mark current node as visited




Visited Nodes: [A, B, D, E, F, C] Unvisited Nodes: []

Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	12	F
D	7	B
E	8	B
F	9	D

Getting the Shortest Path from A to C


- C's previous vertex is F; F's previous vertex is D; D's previous vertex is B; B's previous vertex is A
- Shortest Path from A to C is ABDFC

Get shortest path from A to C




Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	12	F
D	7	B
E	8	B
F	9	D

Dijkstra's Algorithm Example 2




Initialize




N	SD	PN
A	0	
B	∞	
C	∞	
D	∞	
E	∞	

Visit vertex A




N	SD	PN
A	0	
B	3	A
C	1	A
D	∞	
E	∞	

Visit vertex C




N	SD	PN
A	0	
B	2	C
C	1	A
D	∞	
E	5	C

Visit vertex B



N	SD	PN
A	0	
B	2	C
C	1	A
D	5	B
E	3	B


Visit vertex E



N	SD	PN
A	0	
B	2	C
C	1	A
D	5	B
E	3	B

Nothing changes

Visit vertex D




N	SD	PN
A	0	
B	2	C
C	1	A
D	5	B
E	3	B

Nothing changes

Dijkstra's Algorithm Example 3

- Consider vertices in increasing order of distance from s
 - (non-tree vertex with the lowest distTo[] value).
- Add vertex to tree and relax all edges pointing from that vertex.

choose vertex 5
 relax all edges adjacent from 5
 choose vertex 2
 relax all edges adjacent from 2
 choose vertex 3
 relax all edges adjacent from 3
 choose vertex 6
 relax all edges adjacent from 6



choose source vertex 0
 relax all edges adjacent from 0
 choose vertex 1
 relax all edges adjacent from 1


choose vertex 7
 relax all edges adjacent from 7
 choose vertex 4
 relax all edges adjacent from 4

v distTo[]			
→ 0	∞	0	
→ 1	∞	5	
→ 2	∞	17	15 14
→ 3	∞	20	17
→ 4	∞	9	
→ 5	∞	14	13
→ 6	∞	29	26 25
→ 7	∞	8	

v edgeTo[]			
0	-		
1	-	0	
2	-	1	7 5
3	-	1	2
4	-	0	
5	-	7	4
6	-	4	5 2
7	-	0	


Dijkstra's Algorithm Example 4

- Suppose we run Dijkstra's single source shortest-path algorithm on the following edge weighted directed graph with vertex P as the source. In what order do the vertices get included into the set of vertices for which the shortest path distances are finalized?
- ANS: P, Q, R, U, S, T



SD: Shortest Distance

PN: Previous vertex



N	SD	PN
P	0	
Q	∞	
R	∞	
S	∞	
T	∞	
U	∞	

Visit P
→

N	SD	PN
P	0	
Q	1	P
R	∞	
S	6	P
T	7	P
U	∞	

Visit Q
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	5	Q
T	7	P
U	∞	

Visit R
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	R
T	7	P
U	3	R

← Visit U (nothing changes)

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	R
T	7	P
U	3	R

Visit S
(nothing changes)
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	R
T	7	P
U	3	R

Visit T
(nothing changes)
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	R
T	7	P
U	3	R

Finished
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	R
T	7	P
U	3	R

Bellman-Ford Algorithm

- Initialize distance array `distTo[]` for each vertex `v` as `distTo[v] = ∞`, and `distTo[s] = 0` to source vertex `s`.
- Relax all edges $V-1$ times.
 - Can terminate early when all `distTo[]` values have converged
 - The order of edge relaxations affects algorithm efficiency but not correctness.

```
private void relax(DirectedEdge e)
{
    Int u = e.from(), v = e.to();
    if (distTo[v] > distTo[u] + w(u,v))
    {
        distTo[v] = distTo[u] + w(u,v);
        prevNode[v] = u;
    }
}
```

Recall:

Generic algorithm (to compute SPT from `s`)

For each vertex `v`: `distTo[v] = ∞`.

For each vertex `v`: `edgeTo[v] = null`.

`distTo[s] = 0`.

Repeat until done:

- Relax any edge.

Bellman-Ford algorithm

For each vertex `v`: `distTo[v] = ∞`.

For each vertex `v`: `edgeTo[v] = null`.

`distTo[s] = 0`.

Repeat $V-1$ times:

- Relax each edge.

Bellman-Ford Algorithm Proof of Correctness

- Relaxing edges $V-1$ times in the Bellman-Ford algorithm guarantees that the algorithm has explored all possible paths with up to $V-1$ edges, which is the maximum possible number of edges of a shortest path in a graph with V vertices.
- This allows the algorithm to correctly calculate the shortest paths from the source vertex to all other vertices, given that there are no negative-weight cycles.

Bellman-Ford Algorithm with Negative Cycle Detection


- Initialize distance array $\text{distTo}[]$ for each vertex v as $\text{distTo}[v] = \infty$, and $\text{distTo}[s] = 0$ to source vertex s .
- Relax all edges $V-1$ times.
 - Can terminate early when all $\text{distTo}[]$ values have converged
 - The order of edge relaxations affects algorithm efficiency but not correctness. A good heuristic is to follow the Breadth First Search (BFS) order.
- Relax all the edges one more time i.e. the V -th time:
 - Case 1 (Negative cycle exists): if any edge can be further relaxed, i.e., for any edge $u \rightarrow v$, if $\text{distTo}[u] > \text{distTo}[u] + w(u,v)$
 - Case 2 (No Negative cycle) : case 1 fails for all the edges.
- Notes:
 - It can find any negative cycle that is reachable from source vertex s (but not negative cycles that are unreachable from s).
 - If there is a negative cycle that is reachable from source vertex s , then any paths that go through the cycle has distance $-\infty$, since the cost can be reduced by traversing the cycle infinite number of times.

Time Complexity of Bellman-Ford Algorithm

- Time complexity for connected graph:
- Average Case: $O(VE)$
- Worst Case: $O(VE)$
 - If the graph is dense or complete, the value of E becomes $O(V^2)$. So overall time complexity becomes $O(V^3)$

Bellman-Ford Algorithm Example 1

Repeat $V - 1$ times: relax all E edges.



v	distTo[]
0	∞ 0
1	∞ 5
2	∞ 17 14
3	∞ 20 17
4	∞ 9
5	∞ 13
6	∞ 28 20 25
7	∞ 8

v	edgeTo[]
0	-
1	- 0
2	- 1 5
3	- 1 2
4	- 0
5	- 4
6	- 2 5 2
7	- 0

Reverse order of edge relaxations will result in slower convergence

pass 1 pass 2 pass 3 (converged, no further changes, so stop here)

0→1 0→4 0→7 1→2 1→3 1→7 2→3 2→6 3→6 4→5 4→6 4→7 5→2 5→6 7→2 7→5

Order of edge relaxations

Dijkstra's Algorithm vs. Bellman-Ford Algorithm


- Dijkstra's Algorithm:
 - Uses a priority queue to select the next vertex to process.
 - Greedily selects the vertex with the smallest tentative distance to source vertex.
 - Works only on graphs with non-negative edge weights.
- Bellman-Ford Algorithm:
 - Iteratively relaxes all edges $V-1$ times.
 - Does not use a priority queue.
 - Can handle graphs with negative edge weights, and can detect negative cycles.
- Dijkstra's algorithm is faster and more efficient for graphs with non-negative weights; Bellman-Ford Algorithm is more versatile as it can handle negative weights and detect negative cycles, albeit at the cost of lower efficiency.

Dijkstra's Algorithm does not work for Negative Edge Weights

- Dijkstra's Algorithm is greedy and optimal: any vertex that has been visited should have its shortest distance to the source. After visiting A, C, D, we have got D's shortest distance to A is 2, but after visiting D, D's distance to A is updated to -10, which violates the greedy optimal assumption of Dijkstra's Algorithm. Even if you update D's distance to A to -10, its downstream vertex E's distance will not be updated.

N	SD	PN		N	SD	PN		N	SD	PN		N	SD	PN
A	0		Visit A	A	0		Visit C	A	0		Visit D	A	0	
B	∞			B	10	A		B	10	A		B	10	A
C	∞			C	1	A		C	1	A		C	1	A
D	∞			D	∞			D	2	C		D	2	C
E	∞			E	∞			E	∞			E	3	D


N	SD	PN		N	SD	PN
A	0		Visit B	A	0	
B	10	A		B	10	A
C	1	A		C	1	A
D	2	C		D	-10	B
E	3	D		E	3	D



Bellman Ford Algorithm works for Negative Edge Weights

- We run for $V-1=3$ iterations, then run one more iteration with no change. Hence we conclude that The Bellman-Ford algorithm successfully calculated the shortest paths from vertex A to all other vertices. The shortest path from vertex A to vertex D goes through vertex B with a total cost of -10. There are no negative weight cycles.
- Suppose edge update order $A \rightarrow B$, $A \rightarrow C$, $B \rightarrow D$, $C \rightarrow D$, $D \rightarrow E$

N	SD	PN		N	SD	PN		N	SD	PN
A	0			A	0			A	0	
B	∞		Iter 1	B	10	A	Iter 2	B	10	A
C	∞			C	1	A	No change	C	1	A
D	∞			D	-10	B	converged	D	-10	B
E	∞			E	-9	D		E	-9	D



Topological Sort for Shortest Paths in Edge-weighted DAG

- Suppose that a graph is a Directed Acyclic Graph (DAG), i.e., it has no directed cycles. It is easier and faster to find shortest paths than in a general digraph.
- Idea: Consider vertices in topological order. Relax all outgoing edges from that vertex
- *Initialize $dist[] = \{\infty, \infty, \dots\}$ and $dist[s] = 0$ where s is the source vertex.*
- *Create a topological order of all vertices.*
- *For every vertex u in topological order*
 - For every adjacent vertex v of u*
 - if $(dist[v] > dist[u] + weight(u, v))$ //relax edge uv*
 - $dist[v] = dist[u] + weight(u, v)$*
- Time Complexity: Time complexity of topological sort is $O(V+E)$. After finding topological order, the algorithm process all vertices and for every vertex, it runs a loop for all adjacent vertices. Total adjacent vertices in a graph is $O(E)$, so the double for loop has complexity $O(V+E)$. Therefore, overall time complexity is $O(V+E)$.

Shortest Paths in Edge-weighted DAG: Correctness Proof

Proposition. Topological sort algorithm computes SPT in any edge-weighted DAG in time proportional to $E + V$.



edge weights
can be negative!

Pf.

- Each edge $e = v \rightarrow w$ is relaxed exactly once (when v is relaxed),
 - leaving $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.
- Inequality holds until algorithm terminates because:
 - $\text{distTo}[w]$ cannot increase ← $\text{distTo}[\]$ values are monotone decreasing
 - $\text{distTo}[v]$ will not change ← because of topological order, no edge pointing to v will be relaxed after v is relaxed
- Thus, upon termination, shortest-paths optimality conditions hold.


Topological Sort Example 1

- Consider this DAG and a topological order ADBCEF




Initialize

	distTo	edgeTo
A	0	-
B	∞	-
C	∞	-
D	∞	-
E	∞	-
F	∞	-




Visit A

	distTo	edgeTo
A	0	-
B	1	A
C	∞	-
D	1	A
E	∞	-
F	∞	-

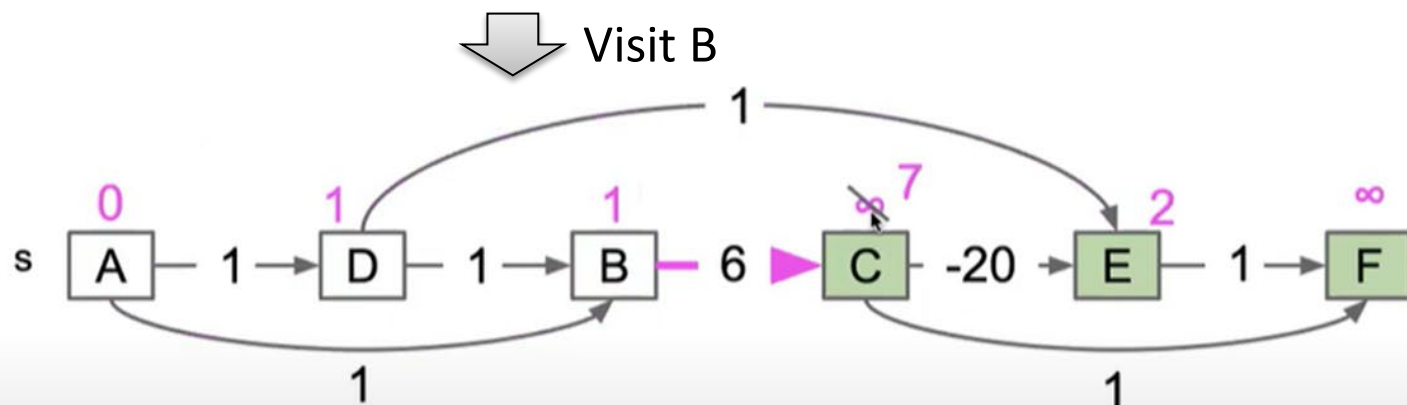


Visit D


	distTo	edgeTo
A	0	-
B	1	A
C	∞	-
D	1	A
E	2	D
F	∞	-




	distTo	edgeTo
A	0	-
B	1	A
C	7	B
D	1	A
E	2	D
F	∞	-



	distTo	edgeTo
A	0	-
B	1	A
C	7	B
D	1	A
E	-13	C
F	8	C




	distTo	edgeTo
A	0	-
B	1	A
C	7	B
D	1	A
E	-13	C
F	-12	E



Topological Sort Example 2

- Consider this DAG and a topological order 01475236



0 1 4 7 5 2 3 6
 ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑

v distTo[]			
0	∞	0	
1	∞	5	
2	∞	17	15 14
3	∞	20	17
4	∞	9	
5	∞	13	
6	∞	29	26 25
7	∞	8	

v edgeTo[]			
0	-		
1	-	0	
2	-	1	7 5
3	-	1	2
4	-	0	
5	-	4	
6	-	4	5 2
7	-	0	

Topological Sort Example 3

N	SD	PN
A	0	
B	∞	
C	∞	
D	∞	
E	∞	

Visit A



N	SD	PN
A	0	
B	10	A
C	1	A
D	∞	
E	∞	

Visit B



N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	∞	

Visit C



N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	∞	

Visit D




N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	-9	D

Visit E



N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	-9	D




Consider topological order ABCDE

Topological Sort Example 3

N	SD	PN
A	0	
B	∞	
C	∞	
D	∞	
E	∞	

Visit A




N	SD	PN
A	0	
B	10	A
C	1	A
D	∞	
E	∞	

Visit C



N	SD	PN
A	0	
B	10	A
C	1	A
D	2	C
E	∞	

Visit B



N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	∞	

Visit D




N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	-9	D

Visit E



N	SD	PN
A	0	
B	10	A
C	1	A
D	-10	B
E	-9	D



Consider topological order ACBDE

Longest Paths in Edge-weighted DAG

Formulate as a shortest paths problem in edge-weighted DAGs.


- Negate all weights.
 - Find shortest paths.
 - Negate weights in result.
- equivalent: reverse sense of equality in relax()

longest paths input

5->4	0.35
4->7	0.37
5->7	0.28
5->1	0.32
4->0	0.38
0->2	0.26
3->7	0.39
1->3	0.29
7->2	0.34
6->2	0.40
3->6	0.52
6->0	0.58
6->4	0.93

shortest paths input

5->4	-0.35
4->7	-0.37
5->7	-0.28
5->1	-0.32
4->0	-0.38
0->2	-0.26
3->7	-0.39
1->3	-0.29
7->2	-0.34
6->2	-0.40
3->6	-0.52
6->0	-0.58
6->4	-0.93



Topological sort algorithm for DAGs works even with negative weights.

For general graphs, the longest paths problem is an unsolved problem (exponential time at best)

Single Source Shortest-paths Algorithms Summary


Algorithm	Restriction	Worst-Case Complexity
Dijkstra (Fibonacci heap)	Undirected or directed graph; no negative weights/cycles	$O(V \log V + E)$
Bellman-Ford	Directed graph with negative weights; undirected graph with no negative weights (since a negative weight edge forms a negative cycle by itself)	$O(EV)$
Topological Sort	Directed Acyclic Graph (DAG) (directed graph, no cycles)	$O(E+V)$

Floyd Warshall Algorithm for all-pairs shortest paths


- The Floyd Warshall Algorithm is an all pair shortest path algorithm unlike Dijkstra and Bellman Ford which are single source shortest path algorithms.
- It works for both the directed and undirected weighted graphs. But, it does not work for the graphs with negative cycles
- It follows Dynamic Programming approach to check every possible path going via every possible vertex in order to calculate shortest distance between every pair of vertices.
- *For $k = 0$ to $n - 1$
 For $i = 0$ to $n - 1$
 For $j = 0$ to $n - 1$
 $dist[i, j] = \min(dist[i, j], dist[i, k] + dist[k, j])$*
- *// Update $dist[i, j]$ if shortcut through vertex k has shorter path*
- *where i = source vertex, j = Destination vertex, k = Intermediate vertex*
- Time Complexity: $O(V^3)$, where V is the number of vertices in the graph and we run three nested loops each of size V

Floyd Warshall Algorithm is Dynamic Programming

- Floyd Warshall Algorithm is a Dynamic Programming based algorithm. It finds all pairs shortest paths using following recursive nature of problem. For every pair (i, j) of source and destination vertices respectively, there are two possible cases. 1) k is not an intermediate vertex in shortest path from i to j . We keep the value of $\text{dist}[i][j]$ as it is. 2) k is an intermediate vertex in shortest path from i to j . We update the value of $\text{dist}[i][j]$ as $\text{dist}[i][k] + \text{dist}[k][j]$. The following figure shows the above optimal substructure property in the all-pairs shortest path problem, which enables the use of dynamic programming.



Floyd Warshall Algorithm Example




- **Step 1:** Initialize the $dist[][]$ matrix using the input graph such that $dist[i][j]$ = weight of edge from i to j , also $dist[i][j] = \infty$ if there is no edge from i to j .

Step1: Initializing Distance[][] using the Input Graph

	A	B	C	D	E
A	0	4	∞	5	∞
B	∞	0	1	∞	6
C	2	∞	0	3	∞
D	∞	∞	1	0	2
E	1	∞	∞	4	0

Example Graph



- **Step 2:** Treat vertex **A** as an intermediate vertex and calculate the $dist[i][j]$ for every $\{i,j\}$ vertex pair using the formula:
- $dist[i][j] = \min(dist[i][j], dist[i][A] + dist[A][j])$

Step 2: Using Node A as the Intermediate node


$$Distance[i][j] = \min (Distance[i][j], Distance[i][A] + Distance[A][j])$$

	A	B	C	D	E
A	0	4	∞	5	∞
B	∞	0	1	∞	6
C	2	∞	0	3	∞
D	∞	∞	1	0	2
E	1	∞	∞	4	0



	A	B	C	D	E
A	0	4	∞	5	∞
B	∞	0	1	∞	6
C	2	6	0	3	∞
D	∞	∞	1	0	2
E	1	5	∞	4	0

Example Graph



- **Step 3:** Treat vertex **B** as an intermediate vertex and calculate the $dist[i][j]$ for every $\{i,j\}$ vertex pair using the formula:

$$dist[i][j] = \text{minimum} (dist[i][j], dist[i][B] + dist[B][j])$$


Step 3: Using Node B as the Intermediate node

$$Distance[i][j] = \min (Distance[i][j], Distance[i][B] + Distance[B][j])$$

	A	B	C	D	E
A	0	4	∞	5	∞
B	∞	0	1	∞	6
C	2	6	0	3	∞
D	∞	∞	1	0	2
E	1	5	∞	4	0

	A	B	C	D	E
A	0	4	5	5	10
B	∞	0	1	∞	6
C	2	6	0	3	12
D	∞	∞	1	0	2
E	1	5	6	4	0

Example Graph



- **Step 4:** Treat vertex **C** as an intermediate vertex and calculate the $dist[i][j]$ for every $\{i,j\}$ vertex pair using the formula:

$$dist[i][j] = \text{minimum} (dist[i][j], dist[i][C] + dist[C][j])$$


Step 4: Using Node C as the Intermediate node

$$Distance[i][j] = \min (Distance[i][j], Distance[i][C] + Distance[C][j])$$

	A	B	C	D	E
A	0	4	5	5	10
B	∞	0	1	∞	6
C	2	6	0	3	12
D	∞	∞	1	0	2
E	1	5	6	4	0

	A	B	C	D	E
A	0	4	5	5	10
B	3	0	1	4	6
C	2	6	0	3	12
D	3	7	1	0	2
E	1	5	6	4	0

Example Graph



- **Step 5:** Treat vertex **D** as an intermediate vertex and calculate the $dist[i][j]$ for every $\{i,j\}$ vertex pair using the formula:

$$dist[i][j] = \text{minimum} (dist[i][j], dist[i][D] + dist[D][j])$$


Step 5: Using Node D as the Intermediate node

$$Distance[i][j] = \min (Distance[i][j], Distance[i][D] + Distance[D][j])$$

	A	B	C	D	E
A	0	4	5	5	10
B	3	0	1	4	6
C	2	6	0	3	12
D	3	7	1	0	2
E	1	5	6	4	0

	A	B	C	D	E
A	0	4	5	5	7
B	3	0	1	4	6
C	2	6	0	3	5
D	3	7	1	0	2
E	1	5	5	4	0

Example Graph



- **Step 6:** Treat vertex **E** as an intermediate vertex and calculate the $dist[i][j]$ for every $\{i,j\}$ vertex pair using the formula:

$$dist[i][j] = \text{minimum} (dist[i][j], dist[i][E] + dist[E][j])$$


Step 6: Using Node E as the Intermediate node

$$Distance[i][j] = \min (Distance[i][j], Distance[i][E] + Distance[E][j])$$

	A	B	C	D	E
A	0	4	5	5	7
B	3	0	1	4	6
C	2	6	0	3	5
D	3	7	1	0	2
E	1	5	5	4	0

	A	B	C	D	E
A	0	4	5	5	7
B	3	0	1	4	6
C	2	6	0	3	5
D	3	7	1	0	2
E	1	5	5	4	0

Example Graph




- **Step 7:** Since all the vertices have been treated as an intermediate vertex, we can now return the updated `dist[][]` matrix as our answer matrix.

Step 7: Return Distance[][] matrix as the result

	A	B	C	D	E
A	0	4	5	5	7
B	3	0	1	4	6
C	2	6	0	3	5
D	3	7	1	0	2
E	1	5	5	4	0

Example Graph







Johnson's Algorithm for all-pairs shortest paths

- Idea: run Dijkstra's Single Source shortest path algorithm with every vertex as the source.
- Dijkstra's algorithm doesn't work for negative weight edge. The idea of Johnson's algorithm is to reweight all edges and make them all positive, then run Dijkstra's algorithm with every vertex as the source.
 - We can run Bellman-Ford algorithm with every vertex as the source without reweighting, since Bellman-Ford algorithm can handle negative edge weights, but the time complexity is much higher than running Dijkstra's algorithm.
- How to transform a given graph into a graph with non-negative weight edges without changing the shortest paths?

Example 1: Increase weight of every edge by a constant?

- True or False: In a weighted graph, assume that the shortest path from source s to destination t is correctly calculated using a shortest path algorithm. If we increase weight of every edge by a constant, the shortest path always remains same.
- False. See the following counterexample. There are 4 edges sa , ab , bt and st with weights 1, 1, 1 and 4 respectively. The shortest path from s to t is $sabt$ with cost 3. If we increase weight of every edge by 1, the shortest path changes to st with cost 5.
- Similarly for negative weight edges. There are 4 edges sa , ab , bt and st with weights -1, -1, -1 and -2 respectively. The shortest path from s to t is $sabt$ with cost -3. If we increase weight of every edge by 1, the shortest path changes to st with cost -1.



Double the original weights?

- True or False: Is the following statement valid about shortest paths? Given a graph, suppose we have calculated shortest path from a source to all other vertices. If we modify the graph such that weights of all edges becomes double of the original weight, then the shortest path remains same, and only the total weight of path changes.
- True. The shortest path remains same. It is like if we change unit of distance from meter to kilo meter, the shortest paths do not change. But this does not make weights positive.


Johnson's algorithm for All-pairs Shortest Paths

1. Let the given graph be G . Add a dummy source vertex d , and add edges with weight 0 from d to all vertices of G . Let the modified graph be G' .
2. Run Bellman-Ford algorithm on G' with d as the source. Let the shortest distances calculated by Bellman-Ford be $h[0], h[1], \dots, h[V-1]$. If we find a negative weight cycle, then return. (We run Bellman-Ford algorithm since it can handle negative edge weights.)
3. Reweight the edges of the original graph. For each edge (u, v) , assign the new weight as $w'(u, v) = w(u, v) + h[u] - h[v]$.
4. Remove the added dummy vertex d , and run Dijkstra's algorithm with every vertex as the source to obtain all-pairs shortest paths. Subtract $h[u] - h[v]$ from length of each shortest path to obtain the lengths of shortest paths in the original graph.

Time complexity: Johnson's algorithm uses both Dijkstra and Bellman-Ford as subroutines. The main steps in the algorithm are Bellman-Ford Algorithm called once and Dijkstra called V times. Time complexity of Bellman Ford is $O(VE)$ and time complexity of Dijkstra is $O(V \log V + E)$. So overall time complexity is $O(V^2 \log V + VE)$.


Johnson's Algorithm: Proof

- The following property is always true since $h[]$ values are the shortest distances from the dummy source code d :
 - $h[v] \leq h[u] + w(u, v)$
- The property states that the shortest distance from u to v must be smaller than or equal to the shortest distance from s to u plus edge weight $w(u, v)$. Because of this inequality, the new weights $w'(u, v) = w(u, v) + h[u] - h[v]$ must be greater than or equal to 0.
- After reweighting, and all weights become non-negative, and all set of paths between any two vertices s and t is increased by the same amount, hence the shortest paths remain the same as the original graph before reweighting.
 - Consider any path between two vertices s and t , the weight of every path is increased by $h[s] - h[t]$, since the added $h[]$ values for all intermediate vertices on the path from s to t cancel each other out.



Johnson's Algorithm Example 2



- We add a dummy source vertex d and add edges with weight 0 from d to all vertices of the original graph.
- We run Bellman-Ford algorithm to calculate the shortest distances from d to all other vertices. The shortest distances from d to 0, 1, 2 and 3 are 0, -5, -1 and 0 respectively, i.e., $h[] = \{0, -5, -1, 0\}$. Once we get these distances, we remove vertex d and reweight each edge uv as: $w'(u, v) = w(u, v) + h[u] - h[v]$.
- Since all weights are greater than or equal to 0 now, we can run Dijkstra's shortest path algorithm with every vertex as the source.



Distances from 4 to 0, 1, 2 and 3 are 0, -5, -1 and 0 respectively.

Johnson's Algorithm: Example 1 Revisited

- After adding a dummy source vertex d and edge reweighting, all edge weights are greater than or equal to 0, so we can run Dijkstra's shortest path algorithm with every vertex as the source.
- For example, the shortest path from s to t remains to be $sabt$, same as the original graph before reweighting. Its length in the original graph is $0 - (h[s] - h[t]) = 0 - (0 - (-3)) = -3$.
- $w'(s, a) = w(s, a) + h[s] - h[a] = -1 + 0 - (-1) = 0$
- $w'(a, b) = w(a, b) + h[a] - h[b] = -1 + -1 - (-2) = 0$
- $w'(b, t) = w(b, t) + h[b] - h[t] = -1 + -2 - (-3) = 0$
- $w'(s, t) = w(s, t) + h[s] - h[t] = -2 + 0 - (-3) = 1$
- $w'(s, a) + w'(a, b) + w'(b, t) = w(s, a) + w(a, b) + w(b, t) + h[s] - h[t] = -1 -1 -1 + 0 - (-3) = 0$
- $w'(s, t) = w(s, t) + h[s] - h[t] = -2 + 0 - (-3) = 1$



Video Tutorials

- Dijkstras Shortest Path Algorithm Explained | With Example | Graph Theory
 - <https://www.youtube.com/watch?v=bZkzH5x0SKU>
 - The following lecture slides are based on this video
- Dijkstra's algorithm in 3 minutes
 - https://www.youtube.com/watch?v=_lHSawdgXpI
- Bellman-Ford in 4 minutes — Theory
 - <https://www.youtube.com/watch?v=9PHkk0UavIM>
- Bellman-Ford in 5 minutes — Step by step example
 - <https://www.youtube.com/watch?v=obWXjtg0L64>
- Shortest Path Algorithms Explained (Dijkstra's & Bellman-Ford)
<https://www.youtube.com/watch?v=AE5I0xACpZs>
- Floyd–Warshall algorithm in 4 minutes
 - <https://www.youtube.com/watch?v=4OQeCuLYj-4>

Tutorials from Geeksforgeeks

- <https://www.geeksforgeeks.org/introduction-to-dijkstras-shortest-path-algorithm/>
- <https://www.geeksforgeeks.org/bellman-ford-algorithm-dp-23/>
- <https://www.geeksforgeeks.org/floyd-warshall-algorithm-dp-16/>
- <https://www.geeksforgeeks.org/johnsons-algorithm/>

Quiz

- Which of the following algorithm can be used to efficiently calculate single source shortest paths in a Directed Acyclic Graph?
 - Dijkstra
 - Bellman-Ford
 - Topological Sort
- ANS: Topological Sort
- Topological Sort has complexity $O(V+E)$, which is the most efficient algorithm among the three

Quiz

- Given a graph where all edges have positive weights, the shortest paths produced by Dijkstra and Bellman Ford algorithm may be different but path weight would always be same.
- ANS: True
- Dijkstra and Bellman-Ford both work fine for a graph with all positive weights, but they are different algorithms and may pick different edges for shortest paths.

Quiz

- Match the following
 - Group A
 - a) Dijkstra's single shortest path algo
 - b) Bellman Ford's single shortest path algo
 - c) Floyd Warshall's all pair shortest path algo
 - Group B
 - p) Dynamic Programming
 - q) Backtracking
 - r) Greedy Algorithm
- Dijkstra is a greedy algorithm where we pick the minimum distant vertex from not yet finalized vertices. Bellman Ford and Floyd Warshall both are Dynamic Programming algorithms where we build the shortest paths in bottom up manner.

Quiz

- Let G be a directed graph whose vertex set is the set of numbers from 1 to 100. There is an edge from a vertex i to a vertex j if either $j = i + 1$ or $j = 3i$. The minimum number of edges in a path in G from vertex 1 to vertex 100 is
- A. 4 B. 7 C. 23 D. 99
- ANS: 7
- The task is to find minimum number of edges in a path in G from vertex 1 to vertex 100 such that we can move to either $i+1$ or $3i$ from a vertex i .
- Since the task is to minimize number of edges, we would prefer to follow $3*i$. Let us follow multiple of 3. $1 \Rightarrow 3 \Rightarrow 9 \Rightarrow 27 \Rightarrow 81$, now we can't follow multiple of 3 anymore. So we will have to follow $i+1$. This solution gives a long path.
- What if we begin from end, and we reduce by 1 if the value is not multiple of 3, else we divide by 3. $100 \Rightarrow 99 \Rightarrow 33 \Rightarrow 11 \Rightarrow 10 \Rightarrow 9 \Rightarrow 3 \Rightarrow 1$
- So we need total 7 edges.